



Message-Passing Programming: MPI

National Tsing Hua University
2025, Fall Semester

Outline

■ MPI Introduction

- History & Evolution

■ Communication Methods

- Synchronous / Asynchronous
- Blocking / Non-Blocking


■ MPI API

- Point-to-Point Communication Routines
- Collective Communication Routines
- Group and Communicator Management Routines


■ MPI-IO

What is MPI

- **MPI = Message Passing Interface**
- A **specification** for the developers and users of message passing libraries
 - By itself, it is an **interface NOT a library**
- Commonly used for **distributed memory system & high-performance computing**
- Goal:
 - **Portable:** Run on different machines or platforms
 - Scalable: Run on million of compute nodes
 - Flexible: Isolate MPI developers from MPI programmers (users)




History and Evolution



- MPI resulted from the efforts of numerous individuals and groups
- Today, MPI implementations are a **combination of MPI-1 and MPI-2**. A few implementations include the full functionality of both
- The MPI Forum is now drafting the **MPI-3 standard**

History and Evolution




Programming Model

■ SPMD: Single Program Multiple Data

- Allow **tasks** to branch or conditionally execute only parts of the program they are designed to execute
- MPI tasks of a parallel program can not be dynamically spawned during run time. (MPI-2 addresses this issue).

■ Distributed memory

- MPI provide method of sending & receiving message



Outline


- MPI Introduction
- Communication Methods
 - Synchronous / Asynchronous
 - Blocking / Non-Blocking
- MPI API
- MPI-IO

Communication Methods


- From the view of the pair of communicated processes
 - **Synchronous communication** --- sending and receiving data occurs simultaneously
 - **Asynchronous communication** --- sending and receiving data occurs non-simultaneously
- From the view of individual function call
 - **Blocking** --- has been used to describe routines that do not return until the transfer is completed
 - **Non-blocking** --- has been used to describe routines that return whether or not the message had been received
- Synchronous vs. blocking:
 - Synchronous comm. commonly implemented by blocking call
 - Synchronous comm. intrinsically performs two action:
Transfer Data & Synchronize Processes

Synchronous/Blocking Message Passing

- **Sender:** wait until the complete message can be accepted by the receiver before sending the message




- **Receiver:** wait until the message it is expecting arrives



Asynchronous/Non-Blocking Message Passing

- How message-passing routines can return before the message transfer has been completed?
 - Generally, a **message buffer** needed between source and destination to hold message
 - Message buffer is a **memory space** at the sender and/or receiver side
 - For send routine, once the **local actions have been completed** and the **message is safely on its way**, the process can continue with subsequent work



Outline

- MPI Introduction
- Communication Methods
- MPI API
 - Getting Start
 - Environment Management Routine
 - Point-to-Point Communication Routines
 - Collective Communication Routines
 - Group and Communicator Management Routines

Getting Start

■ Header file: “mpi.h”

- Required for all programs that make MPI library call

MPI include file

#include “mpi.h”

Declarations, prototypes, etc.

Still executed by
all processes

Program Begins

Serial code

■ MPI calls:

- Format: `rc = MPI_Xxx(parameter, ...)`
- Example: `rc = MPI_Bcast (&buffer,count,datatype,root,comm)`
- Error code: return as “rc”; `rc=MPI_SUCCESS` if successful

Initialize MPI environment

Parallel code begins
`MPI_Init()`

Do work & make message passing calls

Terminate MPI environment

Parallel code ends
`MPI_Finalize()`

■ General MPI program structure:

Serial code


Getting Start

■ Communicators and Groups:

- **Groups** define which **collection of processes** may communicate with each other
- Each group is associated with a **communicator** to perform its communication function calls
- **MPI_COMM_WORLD** is the **pre-defined** communicator for all processors

■ Rank

- An **unique identifier** (task ID) for each process in a **communicator**
- Assigned by the system when the process initializes
- Contiguous and **begin at zero**



Environment Management Routines

■ MPI_Init ()

- Initializes the MPI execution environment
- Must be called before any other MPI functions
- Must be called only once in an MPI program

■ MPI_Finalize ()

- Terminates the MPI execution environment
- No other MPI routines may be called after it

■ MPI_Comm_size (comm, &size)

- Determines the number of processes in the group associated with a communicator

■ MPI_Comm_rank (comm, &rank)

- Determines the rank of the calling process within the communicator
- This rank is often referred to as a task ID

Example

```
#include "mpi.h"
int main (int argc, char *argv[]) {
    int numtasks, rank, rc;
    rc = MPI_Init (&argc,&argv);
    if (rc != MPI_SUCCESS) {
        printf ("Error starting MPI program. Terminating.\n");
        MPI_Abort (MPI_COMM_WORLD, rc);
    }
    MPI_Comm_size (MPI_COMM_WORLD, &numtasks);
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    printf ("Number of tasks= %d My rank= %d\n", numtasks, rank);
    MPI_Finalize ();
}
```

Outline

- MPI Introduction
- Communication Methods
- MPI API
 - Getting Start
 - Environment Management Routine
 - Point-to-Point Communication Routines
 - Collective Communication Routines
 - Group and Communicator Management Routines

Point-to-Point Communication Routines

Blocking send	<code>MPI_Send(buffer,count,type,dest,tag,comm)</code>
Non-blocking send	<code>MPI_Isend(buffer,count,type,dest,tag,comm,request)</code>
Blocking receive	<code>MPI_Recv(buffer,count,type,source,tag,comm,status)</code>
Non-blocking receive	<code>MPI_Irecv(buffer,count,type,source,tag,comm,request)</code>

- buffer: **Address space** that references the data to be sent or received
- type: `MPI_CHAR`, `MPI_SHORT`, `MPI_INT`, `MPI_LONG`, `MPI_DOUBLE`, ...
- count: Indicates the **number of data elements** of a particular type to be sent or received
- comm: indicates the communication context
- source/dest: the **rank** (task ID) of the sender/receiver
- tag: arbitrary non-negative integer **assigned by the programmer** to uniquely **identify** a message. Send and receive operations must match message tags. `MPI_ANY_TAG` is the wild card.
- status: status after operation
- request: used by **non-blocking** send and receive operations

Blocking Example

Blocking send	MPI_Send(buffer,count,type,dest,tag,comm)
Blocking receive	MPI_Recv(buffer,count,type,source,tag,comm,status)

```
MPI_Comm_rank(MPI_COMM_WORLD, &myRank); /* find process rank */
if (myRank == 0) {
    int x=10;
        MPI_Send(&x, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
} else if (myRank == 1) {
    int x;
        MPI_Recv(&x, 1, MPI_INT, 0, MPI_ANY_TAG, MPI_COMM_WORLD, status);
}
```

Non-Blocking Example

Non-Blocking send	MPI_ISend(buffer,count,type,dest,tag,comm,request)
Non-Blocking receive	MPI_IRecv(buffer,count,type,source,tag,comm,request)

```
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);/* find process rank */  
if (myrank == 0) {  
    int x=10;  
    MPI_Isend(&x, 1, MPI_INT, 1, 0, MPI_COMM_WORLD, req1);  
    compute();  
} else if (myrank == 1) {  
    int x;  
    MPI_Irecv(&x, 1, MPI_INT, 0, MPI_ANY_TAG, MPI_COMM_WORLD, req1);  
}  
MPI_Wait(req1, status);
```

- **MPI_Wait()** **blocks** until the operation has actually **completed**
- **MPI_Test()** returns with a flag set indicating whether operation completed at that time.


Outline

- MPI Introduction
- Communication Methods
- **MPI API**
 - Getting Start
 - Environment Management Routine
 - Point-to-Point Communication Routines
 - **Collective Communication Routines**
 - Group and Communicator Management Routines
- MPI-IO

Collective Communication Routines

■ MPI_Barrier (comm)


- Creates a barrier **synchronization** in a group
- Blocks until **all tasks in the group** reach the same MPI_Barrier call



■ MPI_Bcast (&buffer, count, datatype, root, comm)

- Broadcasts (sends) a message from the process with rank "root" to all other processes **in the group**


root=1; count=1;



Collective Communication Routines

- MPI_Scatter (&sendbuf, sendcnt, sendtype, &recvbuf, recvcnt, recvtype, root, comm)
 - Distributes **distinct** messages from a source task to all tasks


- MPI_Gather (&sendbuf, sendcnt, sendtype, &recvbuf, recvcnt, recvtype, root, comm)
 - Gathers **distinct** messages from each task in the group to a single destination task
 - This routine is the **reverse operation** of MPI_Scatter



Collective Communication Routines

- MPI_Reduce (&sendbuf, &recvbuf, count, datatype, op, dest, comm)
 - Applies a **reduction operation** on all tasks in the group and places the result in **one task**

dest=2, count=1; op=MPI_SUM



- Pre-defined Reduction Operations

MPI_MAX	Maximum	MPI_MIN	Minimum
MPI_SUM	Sum	MPI_PROD	Product
MPI_LAND	Logical AND	MPI_BAND	Bit-wise AND
MPI_LOR	Logical OR	MPI_BOR	Bit-wise OR
MPI_LXOR	Logical XOR	MPI_BXOR	Bit-wise XOR


Collective Communication Routines

- **MPI_Allgather (&sendbuf, sendcount, sendtype, &recvbuf, recvcount, recvtype, comm)**

- Concatenation of data **to all tasks**
- This is equivalent to an **MPI_Gather** followed by an **MPI_Bcast**

`sendcnt = recvcnt = 1;`

task0 task1 task2 task3




- **MPI_Allreduce(&sendbuf, &recvbuf, count, datatype, op, comm)**

- Applies a reduction operation and places the result in all tasks
- This is equivalent to an **MPI_Reduce** followed by an **MPI_Bcast**

`count=1; op=MPI_SUM`

task0 task1 task2 task3



Outline

- MPI Introduction
- Communication Methods
- MPI API
 - Getting Start
 - Environment Management Routine
 - Point-to-Point Communication Routines
 - Collective Communication Routines
 - Group and Communicator Management Routines

Group and Communicator Routines

* All these calls are **collective call** that needs to be called by all the processes in the communicator

- Group & Communicator Data Type

- MPI_Group
 - MPI_Comm

- MPI_Comm_group(Comm, &Group)

- Access the group associated with a given communicator

- MPI_Group_incl(Group, size, ranks[], &NewGroup)

- Produce a group by including a subset of members from an existing group
 - all members must be distinct

- MPI_Comm_create(Comm, NewGroup, &NewComm)

- Create a new communicator
 - The new communicator must be a **subset of the original group**

Examples: Divide MPI tasks into two groups

```
int rank, numtasks;  
MPI_Group orig_group, new_group;  
MPI_Comm new_comm  
  
MPI_Init();  
MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
MPI_Comm_size(MPI_COMM_WORLD, &numtasks);  
  
/* Extract the original group handle */  
MPI_Comm_group(MPI_COMM_WORLD, &orig_group);  
  
/* Divide tasks into two distinct groups based upon rank */  
int rank1[4] = {0,1,2,3}; int rank2[4] = {5,6,7,8};  
if (rank < numtasks/2) MPI_Group_incl(orig_group, 4, ranks1, &new_group);  
else MPI_Group_incl(orig_group, 4, ranks2, &new_group);  
  
/* Create new communicator & Broadcast within the new group */  
MPI_Comm_create(MPI_COMM_WORLD, new_group, &new_comm);  
MPI_Barrier(new_comm);  
MPI_Finalize();
```

All MPI tasks must call MPI_Group_incl, but they don't necessarily have to be included in the new group



Outline

■ MPI Introduction

- History & Evolution

■ Communication Methods

- Synchronous / Asynchronous
- Blocking / Non-Blocking


■ MPI API

- Point-to-Point Communication Routines
- Collective Communication Routines
- Group and Communicator Management Routines

■ MPI-IO

Relative Speed of Components in HPC Platform


- An HPC platform's I/O subsystems are typically slow as compared to its other parts
- The I/O gap between memory speed and average disk access stands at roughly 10^{-3}



Concurrent Data Access in a Cluster

We need some magic to make the collection of spinning disks act like a single disk ...

a few hundreds spinning disks





hundreds of thousands of
processors

POSIX File Access Operations

■ POSIX file system call “fopen()”:


- The same file is opened by each processes → multiple file handlers across your MPI processes
- Open the same file with **read permission** is OK
- But can't open with **write permission** together due file system **locking** mechanism → data inconsistency
- To write simultaneously must **create multiple files** (can't take advantage of parallel file system & hard to manage)



MPI-IO File Access Operations

■ MPI-IO call “MPI_File_open()”


- File is opened only once in a collective manner
- MPI library will share and synchronize with each other to use the same file handler
- Can handle both read and write together



MPI-IO Independent/Collective I/O


■ Collective I/O

- Read/write to a shared memory buffer, then issue **ONE file request**
- Reduce #I/O request
→ Good for small I/O
- Require **synchronization**



■ Independent I/O

- Read/write individually
- Prevent synchronization
- One request per process
- Request is **serialized if access the same OST**
- Good for large I/O



MPI-IO API

- **`MPI_File_open(MPI_Comm comm, char *filename, int amode, MPI_Info info, MPI_File *fh)`**
 - Open a file
- **`MPI_File_close(MPI_File *fh)`**
 - Close a file
- **`MPI_File_read/write(MPI_File fh, void *buf, int count, MPI_Datatype datatype, MPI_Status *status)`**
 - Independent read/write using individual file pointer
- **`MPI_File_read/write_all(MPI_File fh, void *buf, int count, MPI_Datatype datatype, MPI_Status *status)`**
 - Collectively read/write using individual file pointer
- **`MPI_File_sync(MPI_File fh)`**
 - Flush all previous writes to the storage device

Reference

- Textbook:

- Parallel Computing Chap2

- MPI Tutorial:

- <https://computing.llnl.gov/tutorials/mpi/>


- MPI API:

- <http://www.mcs.anl.gov/research/projects/mpi/www/www3/>

HW1: Odd-Even Sort

■ Algo:

- comparing & switch in order between all (odd, even)-indexed pairs of adjacent elements in the list
- comparing & switch in order between all (even,odd)-indexed pairs of adjacent elements in the list
- Repeat until the list is sorted



HW1: Odd-Even Sort

■ Sequential code:

```
/* Assumes a is an array of values to be sorted. */
var sorted = false;
while(!sorted) {
    sorted=true;
    for(var i = 1; i < list.length-1; i += 2) {
        if(a[i] > a[i+1]) { swap(a, i, i+1); sorted = false; }
    }
    for(var i = 0; i < list.length-1; i += 2) {
        if(a[i] > a[i+1]) { swap(a, i, i+1); sorted = false; }
    }
}
```

HW1: Odd-Even Sort

■ Parallel Code:

1. For each process with odd rank P , send its number to the process with rank $P-1$.
2. For each process with rank $P-1$, compare its number with the number sent by the process with rank P and send the larger one back to the process with rank P .
3. For each process with even rank Q , send its number to the process with rank $Q-1$.
4. For each process with rank $Q-1$, compare its number with the number sent by the process with rank Q and send the larger one back to the process with rank Q .
5. Repeat 1-4 until the numbers are sorted.