



Shared-Memory Programming: Pthread

National Tsing Hua University
2025, Fall Semester

Outline

- Shared-memory Programming
- Pthread
- Synchronization Problem & Tools

Shared-Memory Programming


- **Definition:** Processes communicate or work together with each other through a shared memory space which can be accessed by all processes
 - Faster & more efficient than message passing

- Many issues as well:

- Synchronization
 - Deadlock
 - Cache coherence

- Programming techniques:


- Parallelizing compiler
 - Unix processes
 - Threads (**Pthread**, Java)



Threads vs. Processes

- **Process (heavyweight process):** complete separate program with its own variables, stack, heap, and everything else.
- **Thread (lightweight process):** share the same memory space for global variables, resources
- In Linux:
 - Threads are created via **clone a process** with a flag to indicate the level of sharing

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.



Why Thread?

- Lower creation/management cost vs. Process

platform	fork()	pthread_create()	speedup
AMD 2.4 GHz Opteron	17.6	1.4	15.6x
IBM 1.5 GHz POWER4	104.5	2.1	49.8x
INTEL 2.4 GHz Xeon	54.9	1.6	34.3x
INTEL 1.4 GHz Itanium2	54.5	2.0	27.3x

- Faster inter-process communication vs. MPI

platform	MPI Shared Memory BW (GB/sec)	Pthreads Worst Case Memory-to-CPU BW (GB/sec)	speedup
AMD 2.4 GHz Opteron	1.2	5.3	4.4x
IBM 1.5 GHz POWER4	2.1	4	1.9x
INTEL 2.4 GHz Xeon	0.3	4.3	14.3x
INTEL 1.4 GHz Itanium2	1.8	6.4	3.6x

Outline

- Shared-memory Programming
- Pthread
 - What is Pthread
 - Pthread Creation
 - Pthread Joining & Detaching
- Synchronization Problem & Tools


What is Pthread?

- Historically, hardware vendors have implemented their own proprietary versions of threads
- **POSIX** (Portable Operating System Interface) standard is specified for portability across Unix-like systems
 - Similar concept as MPI for message passing libraries
- **Pthread** library is the implementation of **POSIX standard** for thread
 - Same relation between MPICH and MPI

Pthread Creation

■ `pthread_create(thread,attr,routine,arg)`

- **thread**: An **unique identifier** (token) for the new thread
- **attr**: It is used to set **thread attributes**. **NULL** for the default values
- **routine**: The routine that the thread will execute once it is created
- **arg**: A **single argument** that may be passed to *routine*



Example

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5


void *PrintHello(void *threadId) {
    int* data = static_cast <int*> (threadId);
    printf("Hello World! It's me, thread #%d!\n", *data);
    pthread_exit(NULL);
}

int main (int argc, char *argv[]) {
    pthread_t threads[NUM_THREADS];
    int tids[NUM_THREADS];
    for(int i=0; i<NUM_THREADS; i++){
        tids[i] = i;
        pthread_create(&threads[i], NULL, PrintHello, (void *)&tids[i]);
    }
    /* Last thing that main() should do */
    pthread_exit(NULL);
}
```

Pthread Joining & Detaching

- `pthread_join(threadId, status)`
 - Blocks until the specified *threadId* thread terminates
 - One way to accomplish synchronization between threads
 - Example: to create a pthread barrier

```
for (int i=0; i<n; i++) pthread_join(thread[i], NULL);
```
- `pthread_detach(threadId)`
 - Once a thread is **detached**, it can **never** be joined
 - Detach a thread could free some system resources



Outline

- Shared-memory Programming
- Pthread
- **Synchronization Problem & Tools**
 - Pthread
 - ◆ Mutually exclusion Lock
 - ◆ Condition variable
 - POSIX Semaphore
 - JAVA Monitor
- Other issues

Synchronization Problem

- The outcome of data content should **NOT** be decided by the **execution order among processes**
- Instructions of individual processes/threads may be **interleaved** in time
- E.g.: Assume variable “**counter**” is shared by processes
- The statement “**counter++**” & “**counter--**” may be implemented in machine language as:

*move ax, counter
add ax, 1
move counter, ax*

*move bx, counter
sub bx, 1
move counter, bx*

Process0
main() {
 :
 counter++;
 :
}

Process1
main() {
 :
 counter--;
 :
}

Instruction Interleaving

- Assume counter is initially 5. One interleaving of statement is:

producer: move ax, counter	→ ax = 5
producer: add ax, 1	→ ax = 6
<i>context switch</i>	
consumer: move bx, counter	→ bx = 5
consumer: sub bx, 1	→ bx = 4
<i>context switch</i>	
producer: move counter, ax	→ counter = 6
<i>context switch</i>	
consumer: move counter, bx	→ counter = 4

- The value of counter may be either 4, 5, or 6
- The ONLY correct result is 5!

Outline

- Shared-memory Programming
- Pthread
- **Synchronization Problem & Tools**
 - Pthread
 - ◆ Mutually exclusion Lock
 - ◆ Condition variable
 - POSIX Semaphore
 - JAVA Monitor
- Other issues

Critical Section & Mutual Exclusion

- **Critical Section** is a piece of code that can only be accessed by one process/thread at a time
- **Mutual exclusion** is the problem to insure only one process/thread can be in a critical section
- E.g.: The design of entry section & exit section provides mutual exclusion for the critical section

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (1);
```

The diagram illustrates the sequence of events in a `do-while` loop for mutual exclusion:

- An arrow labeled "Get entry permission" points from the `entry section` to the `critical section`.
- An arrow labeled "Modify shared data" points from the `critical section` to the `exit section`.
- An arrow labeled "Release entry permission" points from the `exit section` back to the `entry section`.

Locks

- Lock: the simplest mechanism for ensuring mutual exclusion of critical section
 - Spinlock is one of the implementation:

```
while (lock == 1);          /* no operation in while loop */  
lock = 1;                  /* enter critical section */  
  
. . .  
critical section  
  
. . .  
lock = 0;                  /* leave critical section */
```

- Locks are implemented in Pthreads by a special type of variables “mutex”
- **Mutex** is abbreviation of “mutual exclusion”

Pthread Lock/Mutex Routines

- To use mutex, it must be declared as of type `pthread_mutex_t` and initialized with `pthread_mutex_init()`
- A mutex is destroyed with `pthread_mutex_destroy()`
- A critical section can then be protected using `pthread_mutex_lock()` and `pthread_mutex_unlock()`
- Example:

```
#include "pthread.h"
pthread_mutex_t mutex;
pthread_mutex_init (&mutex, NULL);           // specify default
                                                // attribute for the mutex
pthread_mutex_lock(&mutex);                // enter critical section
                                                // Critical Section
                                                // leave critical section
pthread_mutex_unlock(&mutex);
pthread_mutex_destroy(&mutex);
```

Bounded-Buffer Problem


- A pool of n buffers, each capable of holding one item
- Producer:
 - grab an empty buffer
 - place an item into the buffer
 - waits if no empty buffer is available
- Consumer:
 - grab a buffer and retracts the item
 - place the buffer back to the free pool
 - waits if all buffers are empty

Bounded-Buffer Problem

- **Producer** process produces information that is consumed by a **Consumer** process

- Buffer as a circular array with size B

- next free: in
- first available: out
- empty: $in = out$
- full: $(in+1) \% B = out$



- The solution allows at most $(B-1)$ item in the buffer
 - Otherwise, cannot tell the buffer is full or empty


Shared-Memory Solution

```
/*producer*/  
while (1) {  
    while (((in + 1) % BUFFER_SIZE) == out)  
        ; //wait if buffer is full  
    buffer[in] = nextProduced;  
    in = (in + 1) % BUFFER_SIZE;
```


} “in” only modified by producer

```
/*consumer*/  
while (1) {  
    while (in == out); //wait if buffer is empty  
    nextConsumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;
```

} “out” only modified by consumer



```
/* global data structure */  
#define BUFSIZE 10  
item buffer[BUFSIZE];  
int in = out = 0;
```



Using Mutex Lock

```
/*producer*/          /*consumer*/  
while (1) {           while (1) {  
    nextItem = getItem( );  
    while (counter == BUFFER_SIZE) ;  
    buffer[in] = nextItem;  
    in = (in + 1) % BUFFER_SIZE;  
    mutex_lock(mutex);  
    counter++;  
    mutex_unlock(mutex);  
}  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    mutex_lock(mutex);  
    counter--;  
    mutex_unlock(mutex);  
}
```

Condition Variables (CV)

- CV represent some **condition** that a thread can:
 - Wait on, until the condition occurs; or
 - Notify other waiting threads that the condition has occurred
- Three operations on condition variables:
 - **wait()** --- Block until another thread calls **signal()** or **broadcast()** on the CV
 - **signal()** --- Wake up **one thread** waiting on the CV
 - **broadcast()** --- Wake up **all threads** waiting on the CV
- In Pthread, CV **type** is a **pthread_cond_t**
 - Use **pthread_cond_init()** to initialize
 - **pthread_cond_wait (&theCV, &somelock)**
 - **pthread_cond_signal (&theCV)**
 - **pthread_cond_broadcast (&theCV)**

Using Condition Variable

■ Example:

- A threads is designed to take action when $x=0$
- Another thread is responsible for decrementing the counter

```
pthread_cond_t cond;  
pthread_cond_init (cond, NULL);
```

```
pthread_mutex_t mutex;  
pthread_mutex_init (mutex, NULL);
```

```
action() {  
    pthread_mutex_lock (&mutex)  
    if (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

- All condition variable operation **MUST** be performed while a mutex is **locked!!!**

Why is the lock necessary???

Using Condition Variable

```
pthread_cond_t cond;  
pthread_cond_init (cond, NULL);
```

```
action() {  
    pthread_mutex_lock (&mutex)  
    if (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
pthread_mutex_t mutex;  
pthread_mutex_init (mutex, NULL);
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

Because event counter “x” is a **SHARED** variable

- If no lock on thread action()...
 - Wait after **any thread** (i.e. not counter) sets “x” to 0
- If no lock on thread counter()...
 - No guarantee that decrement and test of “x” is **atomic**
- Requiring CV operations to be done while holding a lock
prevents a lot of common programming mistakes

Using Condition Variable

```
action() {  
    →pthread_mutex_lock (&mutex)  
    if (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

■ What really happens...

1. Lock mutex

Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    if (x != 0)  
        → pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    → pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

■ What really happens...

1. Lock mutex
2. Wait()
 1. Put the thread into sleep &
releases the lock

1. Lock mutex

Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    if (x != 0)  
        → pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        → pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

■ What really happens...

1. Lock mutex
2. Wait()
 1. Put the thread into sleep &
releases the lock
 1. Waked up, but the **thread is locked**

1. Lock mutex
2. Signal()

Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    if (x != 0)  
        → pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        → pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

■ What really happens...

1. Lock mutex
2. Wait()
 1. Put the thread into **sleep & releases the lock**
 1. Waked up, but the **thread is locked**
 2. **Re-acquire lock** and resume execution
1. Lock mutex
2. Signal()
3. Releases the lock

Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    if (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

■ What really happens...

1. Lock mutex
2. Wait()
 1. Put the thread into **sleep & releases the lock**
 1. Waked up, but the **thread is locked**
 2. **Re-acquire lock** and resume execution
 3. Release the lock
1. Lock mutex
2. Signal()
3. Releases the lock

Using Condition Variable

```
action() {  
    pthread_mutex_lock (&mutex)  
    if (x != 0)  
        pthread_cond_wait (cond, mutex);  
    pthread_mutex_unlock (&mutex);  
    take_action();  
}
```

```
counter() {  
    pthread_mutex_lock (&mutex)  
    x--;  
    if (x==0)  
        pthread_cond_signal (cond);  
    pthread_mutex_unlock (&mutex);  
}
```

■ What really happens...

1. Lock mutex
 2. Wait()
 1. Put the thread into sleep & **releases the lock**
 1. Waked up, but the **thread is locked**
 2. **Re-acquire lock** and resume execution
 3. Release the lock
-
1. Lock mutex
2. Signal()
3. Releases the lock
- Another reason why condition variable op. MUST within mutex lock

Thread Pools

- Create a number of threads in a pool where they await work
- Advantages
 - Usually slightly faster to service a request with an existing thread than create a new thread
 - Allows the number of threads in the application(s) to be bound to the size of the pool
- # of threads: # of CPUs, expected # of requests, amount of physical memory

ThreadPool Implementation

Task structure

```
typedef struct {
    void (*function)(void *);
    void *argument;
} threadpool_task_t;
```

Allocate thread and task queue

```
/* Allocate thread and task queue */
pool->threads = (pthread_t *) malloc(sizeof(pthread_t) * thread_count);
pool->queue = (threadpool_task_t *) malloc(sizeof(threadpool_task_t) * queue_size);
```

Threadpool structure

```
struct threadpool_t {
    pthread_mutex_t lock;
    pthread_cond_t notify;
    pthread_t *threads;
    threadpool_task_t *queue;
    int thread_count;
    int queue_size;
    int head;
    int tail;
    int count;
    int shutdown;
    int started;
};
```

ThreadPool Implementation

```
static void *threadpool_thread(void *threadpool) ← thread handler  
{  
    threadpool_t *pool = (threadpool_t *)threadpool;  
    threadpool_task_t task;  
  
    for(;;) {  
        /* Lock must be taken to wait on conditional variable */  
        pthread_mutex_lock(&(pool->lock));  
  
        /* Wait on condition variable, check for spurious wakeups.  
         * When returning from pthread_cond_wait(), we own the lock. */  
        while((pool->count == 0) && (!pool->shutdown)) {  
            pthread_cond_wait(&(pool->notify), &(pool->lock));  
        }  
    }  
}
```

ThreadPool Implementation

```
/* Grab our task */

task.function = pool->queue[pool->head].function;
task.argument = pool->queue[pool->head].argument;

pool->head += 1;

pool->head = (pool->head == pool->queue_size) ? 0 : pool->head;

pool->count -= 1;

/* Unlock */

pthread_mutex_unlock(&(pool->lock));

/* Get to work */

(*(task.function))(task.argument);

}
```

Semaphore

- A tool to generalize the synchronization problem
 - Deadlock may occur if not use appropriately !
- More specifically...
 - a record of how many units of a particular resource are available
 - ◆ If #record = 1 ➔ binary semaphore, mutex lock
 - ◆ If #record > 1 ➔ counting semaphore
 - accessed only through 2 *atomic* ops: **wait** & **signal**
- Spinlock implementation:
 - Semaphore is an integer variable

<pre>wait (S) { while (S <= 0) ; S--; }</pre>	<pre>signal (S) { S++; }</pre>
--	--

POSIX Semaphore

- Semaphore is part of **POSIX standard** BUT it is **not belonged to Pthread**

- It can be used with or **without** thread

- POSIX Semaphore routines:

- **sem_init(sem_t *sem, int pshared, unsigned int value)**
 - **sem_wait(sem_t *sem)**
 - **sem_post(sem_t *sem)**
 - **sem_getvalue(sem_t *sem, int *valptr)**
 - **sem_destory(sem_t *sem)**
- Initial value of the semaphore
- Current value of the semaphore

- Example:

```
#include <semaphore.h>
sem_t sem;
sem_init(&sem);
sem_wait(&sem);
// critical section
sem_post(&sem);
sem_destory(&sem);
```

Semaphore Drawback

- Although semaphores provide a convenient and effective synchronization mechanism, its correctness is depending on the programmer
 - All processes access a shared data object must execute `wait()` and `signal()` in the right order and right place
 - This may not be true because honest programming error or uncooperative programmer


Monitor

- High-level synchronization construct that allows the safe sharing of an abstract data type among concurrent processes

Syntax

```
monitor monitor-name {  
    // shared variable declarations  
    procedure body P1 (...) {  
        ...  
    }  
    procedure body P2 (...) {  
        ...  
    }  
    procedure body Pn (...) {  
        ...  
    }  
    initialization code {  
    }  
}
```

Schematic View



Synchronized Tools in JAVA

■ Synchronized Methods (Monitor)

- Synchronized method uses the method receiver as a lock
- Two invocations of synchronized **methods cannot interleave on the same object**
- When one thread is executing a synchronized method for an object, all other threads that invoke synchronized methods for the same object block until the first thread exists the object

```
public class SynchronizedCounter {  
    private int c = 0;  
  
    public synchronized void increment() { c++; }  
  
    public synchronized void decrement() { c--; }  
  
    public synchronized int value() { return c; }  
}
```

Synchronized Tools in JAVA

■ Synchronized Statement (Mutex Lock)

- Synchronized blocks uses the **expression** as a lock
- A synchronized Statement can only be executed once the thread has obtained a **lock** for the object or the class that has been referred to in the statement
- useful for improving concurrency **with fine-grained**

```
public void run()
{
    synchronized(p1)
    {
        int i = 10; // statement without locking requirement
        p1.display(s1);
    }
}
```

The Big Picture

- Getting synchronization right **is hard!**
- How to pick between locks, semaphores, convars, monitors???
- **Locks** are very simple for many cases
 - But may not be the most efficient solution
- **Condition variables** allow threads to sleep while holding a lock
 - Be aware whether they use Mesa or Hoare semantics
- **Semaphores** provide general functionality
 - But also make it really easy to mess up or cause deadlock
- **Monitors** are a “pattern” for using locks and condition variables

Reference

- Textbook:
 - Parallel Computing Chap8
- Pthread Tutorial
 - <https://computing.llnl.gov/tutorials/pthreads/>
- Synchronization Tools:
 - <http://www.eecs.harvard.edu/~mdw/course/cs61/mediawiki/images/7/7e/Lectures-semaphores.pdf>
- Pthread API:
 - <http://www.yolinux.com/TUTORIALS/LinuxTutorialPosixThreads.html>
- JAVA Synchronized methods
 - <http://docs.oracle.com/javase/tutorial/essential/concurrency/syncmeth.html>