



# Performance (Memory) Optimization

National Tsing Hua University  
2025, Fall Semester

# Communication vs Computation

## ■ Peak performance for Kepler

- The peak processing performance is 3935 Gflops.
- The bandwidth is 250GB/s, which equals to 63G floating point data per second.
- The ratio is about 60 times

## ■ Instruction execution

- Each computation **instruction takes 1~4 cycles**
- Each load/store instruction for **global memory access takes 400~800 cycles**
- Memory access to **shared memory can be 1~20 cycles**
- The ratio is about 100 times

# Data Pre-fetch and Reuse

- GPU has faster memory spaces (**but smaller**)
  - Shared memory / L1 cache
  - Register file
- Solution:
  - Hardware: prefetch data to shared memory or registers for later computation (hardware)
  - Software/Programmer: **minimize memory usage & reuse the data in shared memory or registers** as many times as possible

# Outline

## ■ Host memory

- Pined memory
- Asynchronous computation & data transfer
- Streams

## ■ Global/Local memory

- Memory coalescing
- Tiled algorithm

## ■ Shared memory

- Bank conflicts avoidance
- Memory padding

## ■ Address linearization

# Outline

## ■ Host memory

- Pined memory
- Asynchronous computation & data transfer
- Streams

## ■ Global/Local memory

- Memory coalescing
- Tiled algorithm

## ■ Shared memory

- Bank conflicts avoidance
- Memory padding

## ■ Address linearization

# 1. Page-Locked Data Transfers

- `cudaMallocHost()` allows allocation of page-locked (“pinned”) host memory

```
cudaMalloc ( &dev1, size ) ;  
cudaMallocHost( &host1, size ) ;  
...  
cudaMemcpy ( dev1, host1, size, H2D ) ;
```

- Enables highest `cudaMemcpy` performance
- Use with caution!!
  - Allocating too much page-locked memory can reduce overall system (host) performance

## 2. Overlap CPU & GPU Computations

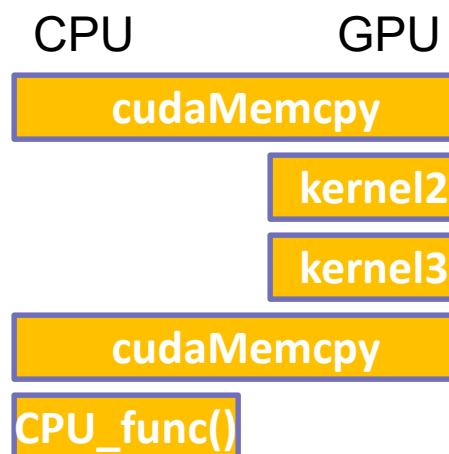
- To facilitate concurrent execution between **host** and **device**, some function calls are **asynchronous**:
  - Control is returned to the host thread before the device has completed the requested task.
- Asynchronous functions:
  - Kernel launches
  - Asynchronous memory copy and set options:  
**cudaMemcpyAsync**, **cudaMemsetAsync**
  - **cudaMemcpy** within the same device
  - H2D cudaMemcpy of **64kB or less**

# Synchronous Computation

```
cudaMalloc ( &dev1, size ) ;  
double* host1 = (double*) malloc ( &host1, size ) ;  
...  
// cudaMemcpy blocks until copy is completed  
cudaMemcpy ( dev1, host1, size, H2D ) ;  
// two kernels are serialized and executed on device  
kernel2 <<< grid, block>>> ( ..., dev2, ... ) ;  
kernel3 <<< grid, block>>> ( ..., dev3, ... ) ;  
// cudaMemcpy starts after kernels finish  
// and blocks until copy is completed  
cudaMemcpy ( host4, dev4, size, D2H ) ;  
CPU_func();  
...
```

Kernels from a single thread are serialized

- CPU and GPU are synchronized due to cudaMemcpy
- Kernel functions from the same process (default stream) are always serialized, and cannot be overlapped on GPU



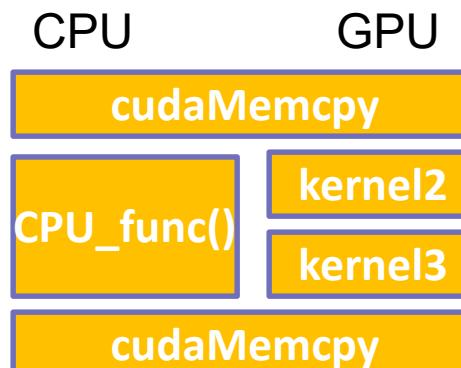
# Asynchronous Computation

```
cudaMalloc(&dev1, size) ;  
double* host1=(double*) malloc (&host1, size);  
...  
cudaMemcpy (dev1, host1, size, H2D) ;  
kernel2 <<< grid, block >>> ( ..., dev1, ... );  
kernel3 <<< grid, block >>> ( ..., dev1, ... );
```

```
CPU_method () ;
```

```
cudaMemcpy ( host1, dev1, size, D2H ) ;  
...
```

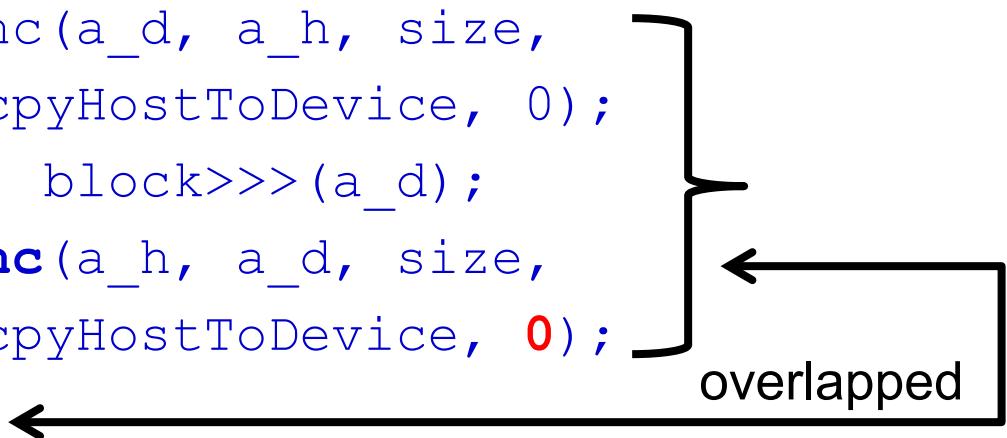
CPU & GPU  
overlapped



# Asynchronous Data Transfers

- Asynchronous host-device memory copy returns control immediately to CPU
  - `cudaMemcpyAsync(dst, src, size, dir, stream);`
  - **requires pinned host memory** (allocated by “`cudaMallocHost()`”)
- Overlap CPU computation with data transfer
  - **0 = default stream**

```
cudaMemcpyAsync(a_d, a_h, size,  
                cudaMemcpyHostToDevice, 0) ;  
  
kernel<<<grid, block>>>(a_d) ;  
  
cudaMemcpyAsync(a_h, a_d, size,  
                cudaMemcpyHostToDevice, 0) ; }  
  
CPU_method();
```

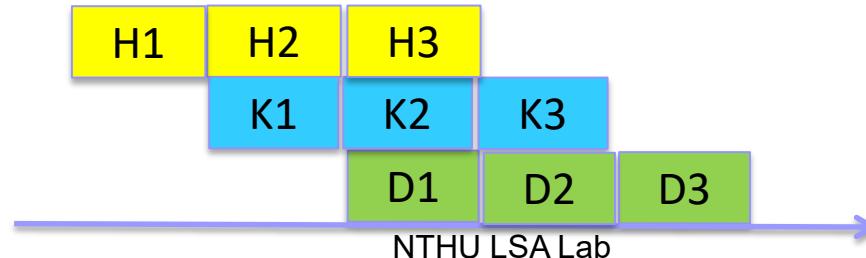


# 3. CUDA Streams

- CUDA Stream is a technique to **overlap the execution of a kernel, and hide data transfer delay from computations**
  - Operations in **different streams** can be interleaved and, when possible, they can even run concurrently (**except steam 0**)
  - **Operations in the same stream are still serialized and executed in order**
- Consider a kernel process a huge dataset
  - Without stream, the kernel computation can only start after the dataset is transferred



- With stream, we can partition the dataset, assign each partition to a stream, and execute them in a **pipeline**



# CUDA Streams

- kernel launch
  - `kernel<<<grid,block,0,stream-id>>> /*...*/;`
- Stream-id must be allocated and destroyed
  - `cudaStream_t *stream;`
  - `cudaStreamCreate(&stream);`
  - `cudaStreamDestroy(stream);`
- Memory copy can be either synchronous or asynchronous. **But synchronous memcpy prevents streams from running in parallel**
- If asynchronous copy is used, host memory must be pinned

# CUDA Streams

```
cudaStream_t stream[2];
cudaStreamCreate(&stream[0]);
cudaStreamCreate(&stream[1]);          pined(page locked mem)
cudaMallocHost(&hostPtr, 2 * size);
for (int i = 0; i < 2; ++i) {
    cudaMemcpyAsync(/*...*/, cudaMemcpyHostToDevice, stream[i]);
    kernel<<<100, 512, 0, stream[i]>>> /*...*/;
    cudaMemcpyAsync(/*...*/, cudaMemcpyDeviceToHost, stream[i]);
}
cudaStreamDestroy(stream[0]);
cudaStreamDestroy(stream[1]);
```

# Stream based Synchronization

- `cudaStreamSynchronize(stream-id)`
  - Blocks host until **all CUDA calls in stream stream-id complete**
- `cudaEventRecord(event, stream-id)`
  - Insert '**events**' in streams
  - Event is recorded when GPU reaches it in a stream
- `cudaEventSynchronize(event)`
  - Blocks **CPU thread** until event is recorded
- `cudaStreamWaitEvent(steam-id, event, 0)`
  - Block a **GPU stream** until event reports completion

# Example: Explicit Sync between Streams

```
cudaEvent_t event;  
cudaEventCreate (&event); // create event  
// 1) H2D copy of new input  
cudaMemcpyAsync ( d_in, in, size, H2D, stream1 );  
cudaEventRecord (event, stream1); // record event  
// 2) D2H copy of previous result  
cudaMemcpyAsync ( out, d_out, size, D2H, stream2 );  
// wait for event in stream1  
cudaStreamWaitEvent ( stream2, event );  
// 3) must wait for 1 and 2  
kernel <<< , , , stream2 >>> ( d_in, d_out );  
asynchronousCPUMethod ( ... ) // Async GPU method
```

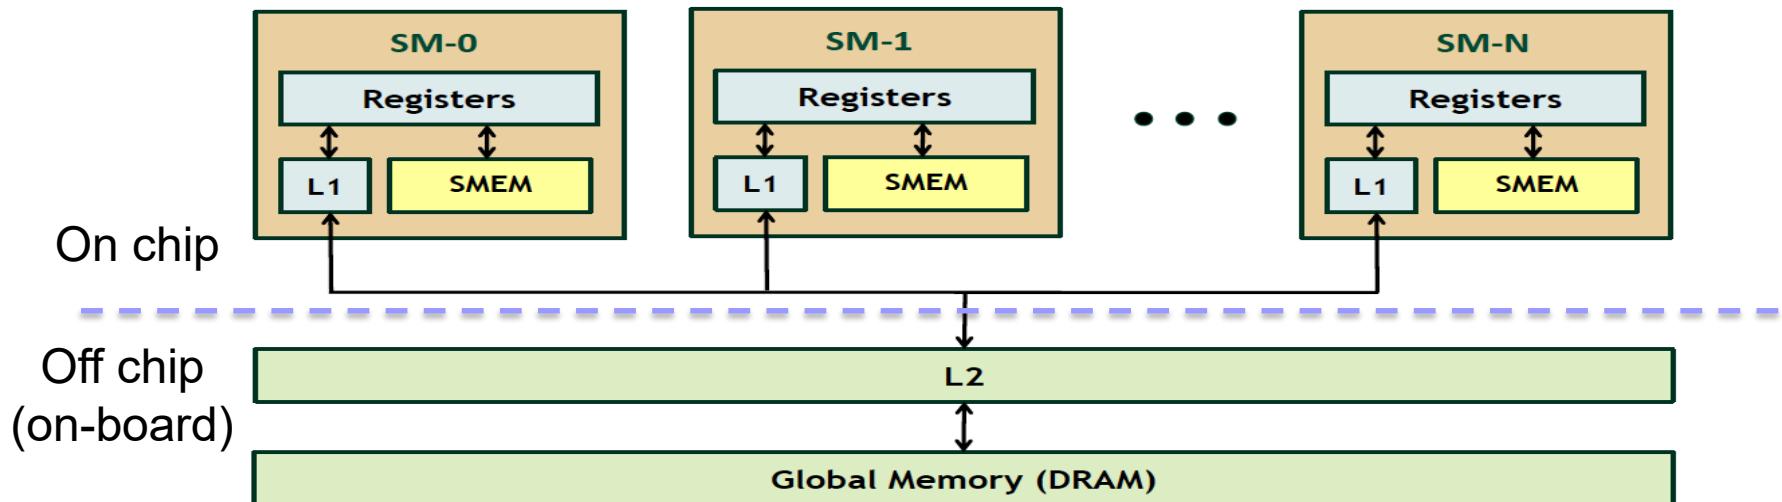


# Outline

- Host memory
  - Pined memory
  - Asynchronous computation & data transfer
  - Streams
- Global/Local memory
  - Memory coalescing
  - Tiled algorithm
- Shared memory
  - Bank conflicts avoidance
  - Memory padding
- Address linearization

# Local Memory Cache

- L1 & L2 are used to cache local memory contents
  - L1: On chip memory. Same as share memory
    - ◆ Programmers can decide the ratio of shared memory and L1 cache
  - L2: Off chip memory Cache. Same as global memory

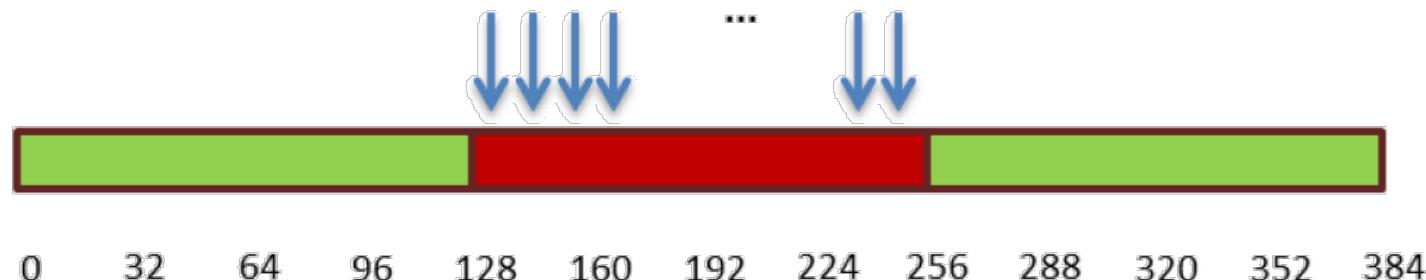


# Coalesced Memory Access

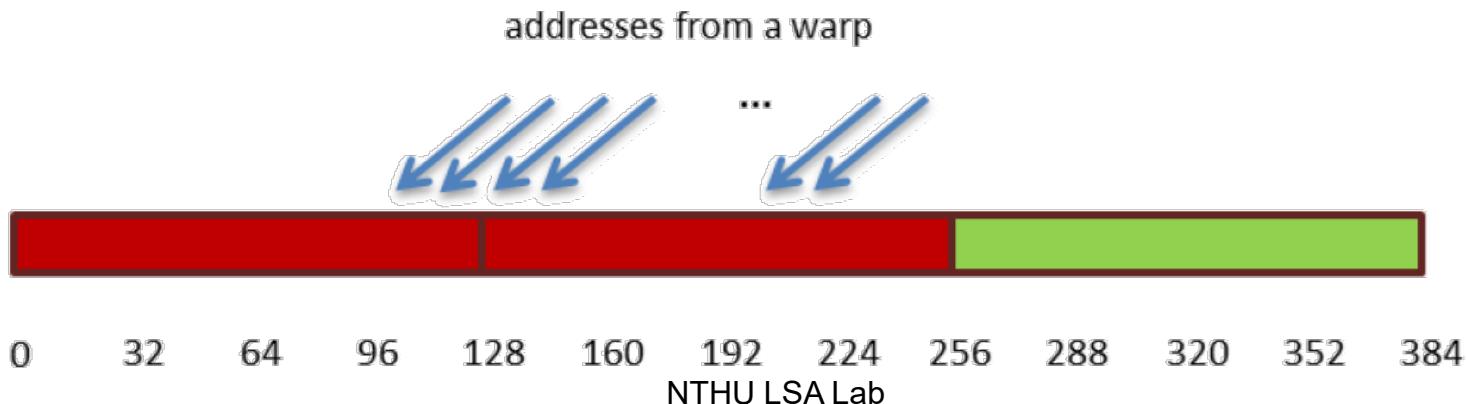
- Accessing data in the global memory is critical to the performance of a CUDA application
  - DRAM is slow comparing to other on-chip memory
- Recall that all threads in a warp execute the same instruction
  - When all threads in a warp execute a load instruction, the hardware detects whether the threads access consecutive memory locations
  - In this favorable case, the hardware coalesces all memory accesses into a consolidated access (single transaction) to consecutive DRAM locations (off-chip memory)

# Coalesced Memory Access

- Coalesced access addresses from a warp

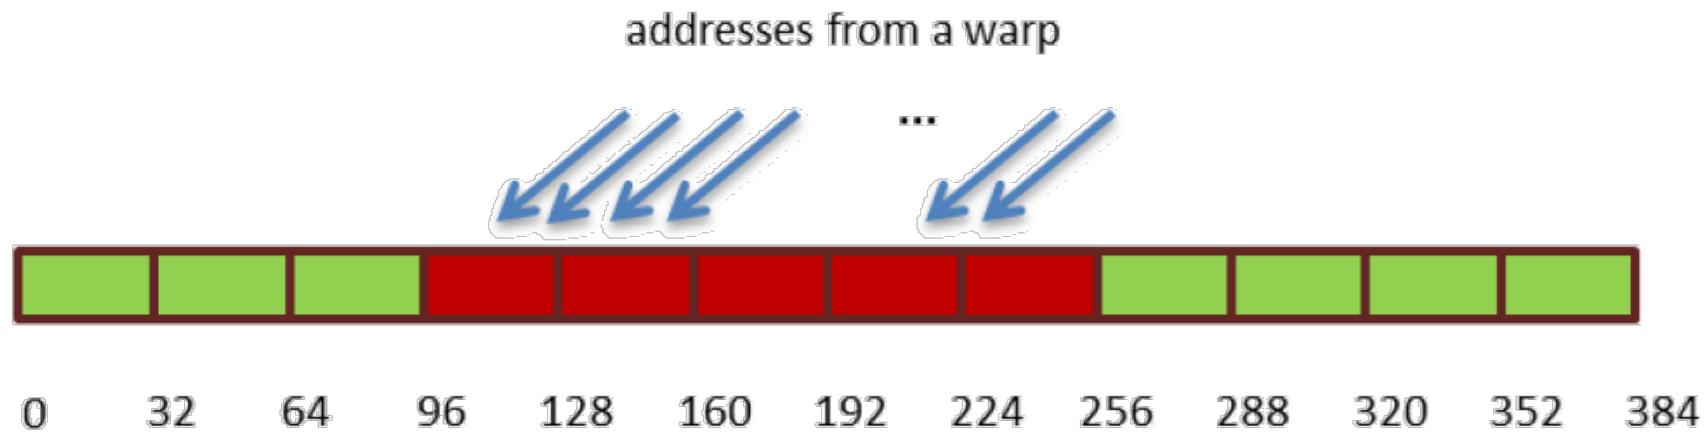


- Unaligned sequential addresses that fit into two 128-byte L1-cache lines



# Misaligned Access Without Caching

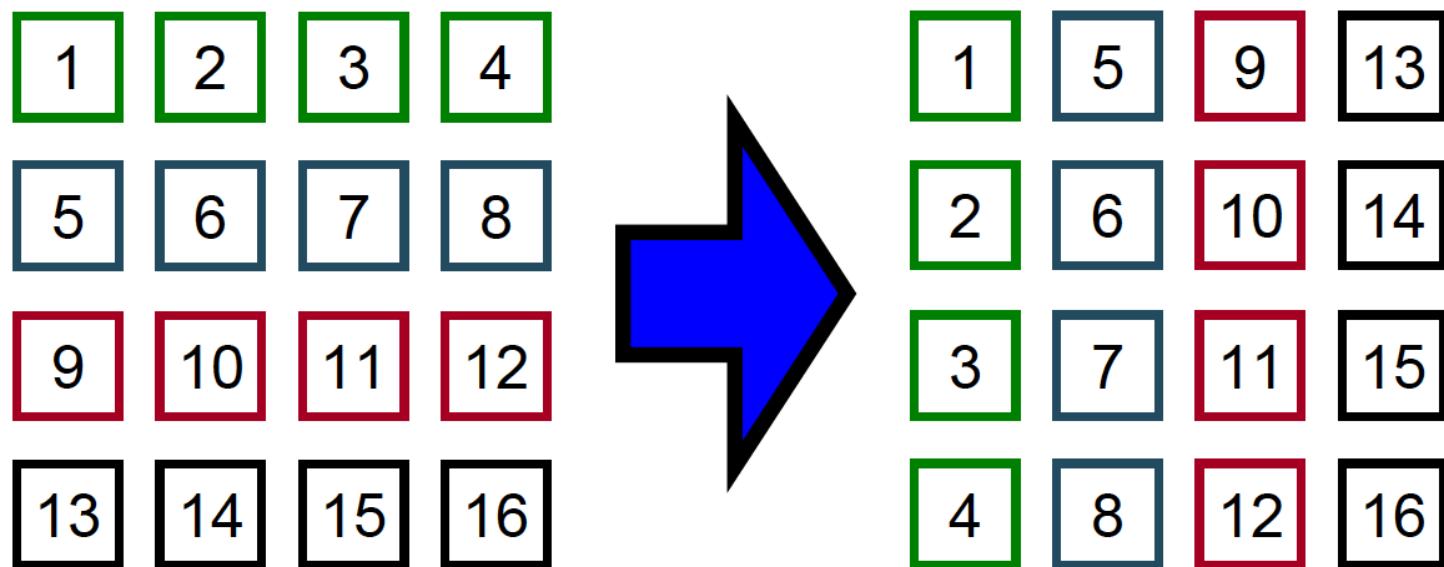
- Misaligned sequential addresses that fall within five **32-byte L2 cache segments**
  - No extra data reading



- Sometimes, it will be faster than (L1) cached memory access
  - If data are not reused

# Example: Matrix Transpose

- SDK Sample (“transpose”)
- Illustrates coalescing using shared memory
  - Speedups for even small matrices



# Uncoalesced Transpose

Reads input from GMEM



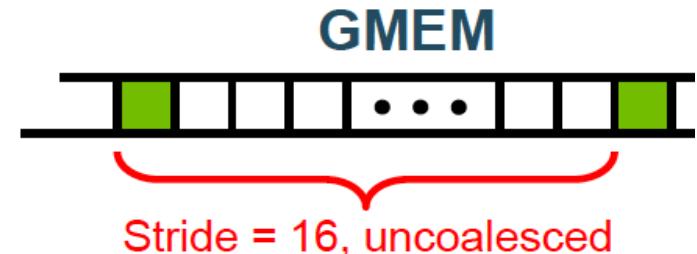
Write output to GMEM



$$B[i,j] = A[j,i]$$



ide = 1, coalesced

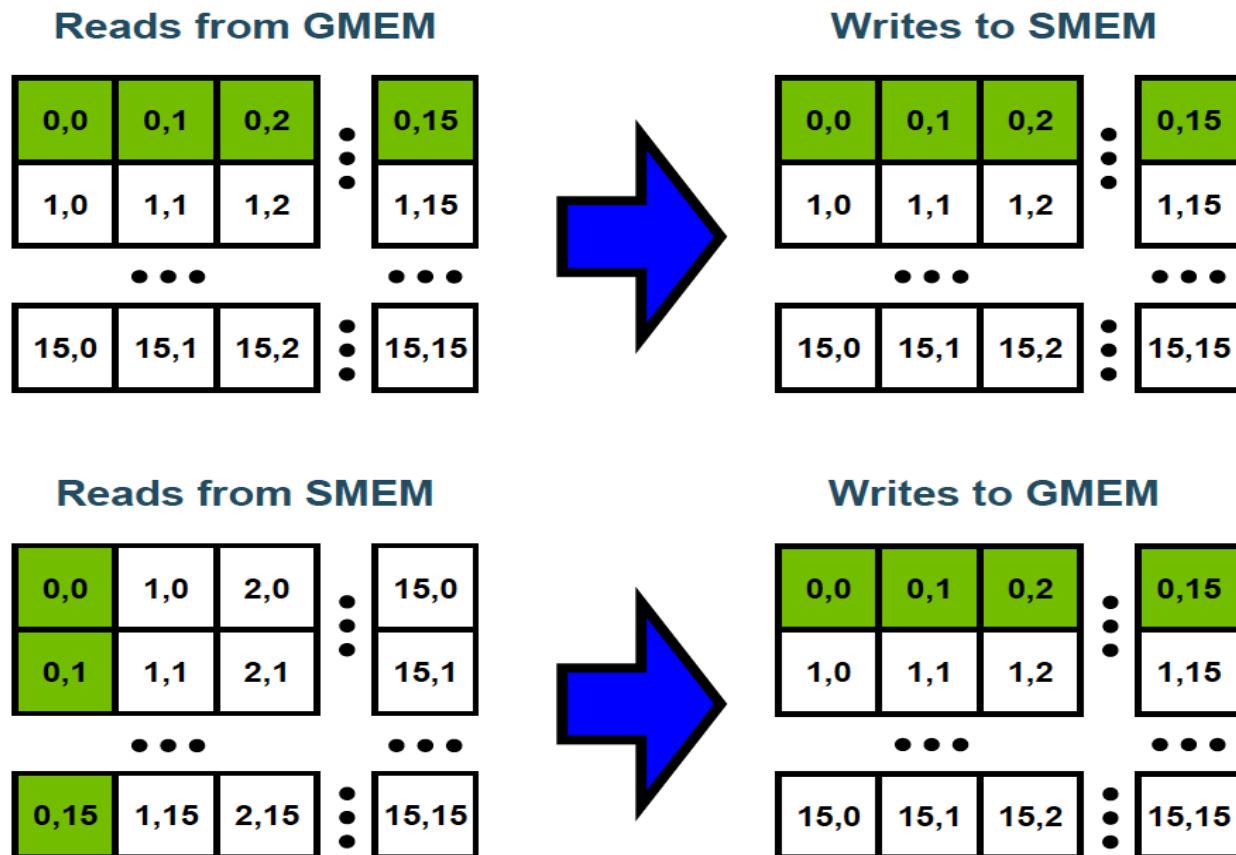


Stride = 16, uncoalesced

# Coalesced Transpose

- Coalescing through shared memory
- Make both read & write become continuous for global memory

share \_\_S[];  
 $S[i,j] = A[i,j];$   
 $B[i,j] = S[j,i];$



# Outline

- Host memory
  - Pined memory
  - Asynchronous computation & data transfer
  - Streams
- Global/Local memory
  - Memory coalescing
  - Tiled algorithm
- Shared memory
  - Bank conflicts avoidance
  - Memory padding
- Address linearization

# Example: Matrix Multiply

## ■ Compute $C = A \times B$ , where $A, B, C$ are $N$ by $N$ matrices

```
For i = 1:N
```

Let each thread compute one element  $C[i][j]$

```
For j = 1:N
```

```
For k = 1:N
```

```
C[i][j] += A[i][k] * B[k][j]
```

## ■ Compute to Global Memory Access (CGMA) ratio

- Compute = 1 multiplication + 1 addition; Memory access = 2
- ➔ CGMA = 1

## ■ K20x (Kepler)

- Compute = 3950 GFLOPs; Global memory BW = 250GB/s
- Compute / Comm. =  $3950 \times 4 / 250 \approx 64$
- ➔ CGMA must increase to 64! Floating point takes 4 bytes

# Load Everything to Shared Memory

- Share memory is 100 times faster than global memory
- If  $N^2$  threads are used:
  - Each thread only needs to loads 2 element, and does  $2N$  computations
  - CGMA = N (When  $N > 64$ , memory access will not be the bottleneck anymore)

```
For i = 1:N
    For j = 1:N
        For k = 1:N
            C[i][j] += A[i][k] * B[k][j]
```

- But shared memory is small
  - The data needs to be stored is  $3N^2$  integers or floats
  - If  $N=1024$ , size = 12MB (i.e.,  $3*1,024*1,024*4$ )

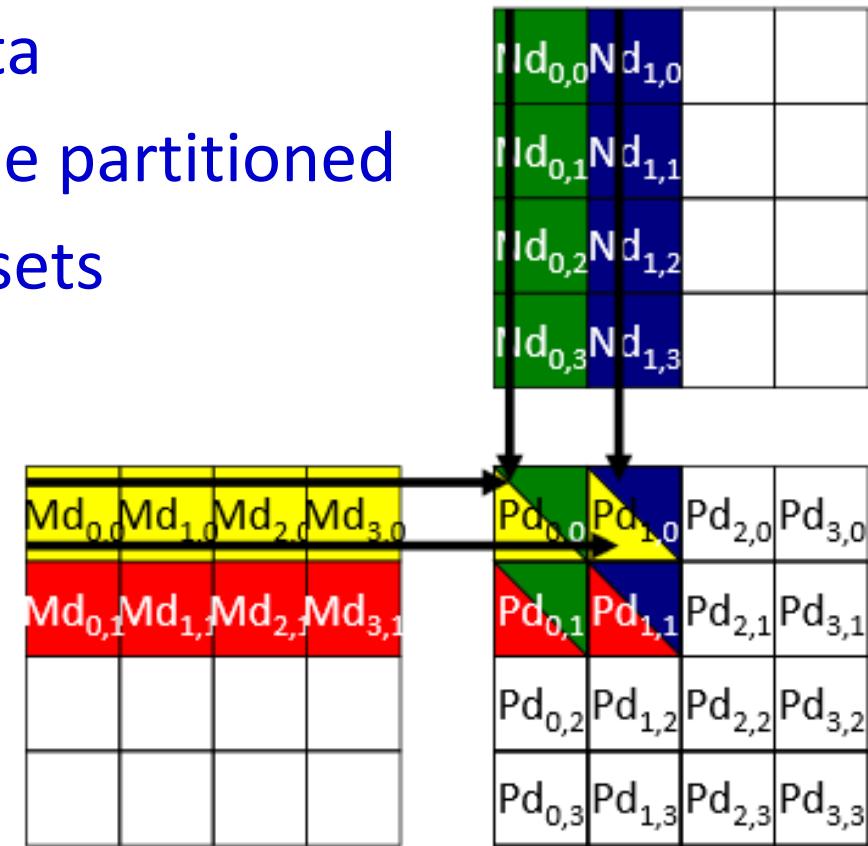
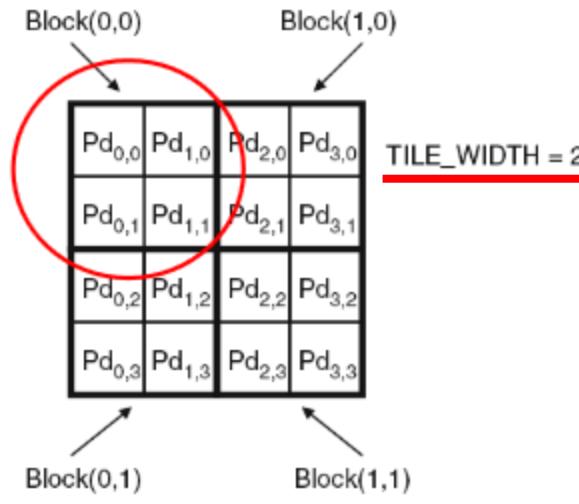
# Load Everything to Shared Memory

- `Matrix_Mul<<<1, N, 2*N*N>>>(A, B, C, N);`
  - The third parameter is the size of shared memory.

```
extern __shared__ int S[];
inline int Addr(int matrixIdx, int i, int j, int N)  {
    return (N*N*matrixIdx + i*N+ j);
}
__global__ void Matrix_Mul(int* A, int* B,int* C, int* N) {
    int i = threadIdx.x;
    int j = threadIdx.y;
    //move data to shared memory
    S[Addr(0, i, j, N)]=A[Addr(0, i, j, N)];
    S[Addr(1, i, j, N)]=B[Addr(0, i, j, N)];
    __syncthreads();
    // do computation
    for(int k=0; k<*N; k++)
        C[Addr(1, i, j, N)]=S[Addr(0, j, k, N)]*S[Addr(0, k, j, N)];
}
```

# Block(Tiled) Algorithm

- Break up the execution of the kernel into phases so that the data accesses in each phase is focused on one subset (tile) of data
- Not all problems can be partitioned into independent subsets



# Block(Tiled) Algorithm

## ■ Rewrite for-loop by TILE\_WIDTH

```
For i' = 1:N step TILE_WIDTH
    For j' = 1:N step TILE_WIDTH
        For k' = 1:N step TILE_WIDTH
            For i = i': i'+TILE_WIDTH - 1
                For j = j': j'+TILE_WIDTH - 1
                    For k = k': k'+TILE_WIDTH - 1
                        C[i][j] += A[i][k]*B[k][j]
```

$$\begin{aligned} \text{Total required data accesses} \\ = 2 \times (\text{TILE\_WIDTH})^2 \end{aligned}$$

$$\text{Total computing} = 2 \times (\text{TILE\_WIDTH})^3$$

- We can find a small enough TILE\_WIDTH, such that all the values needed by  $C[i][j]$  are in shared memory
  - Every data is re-used TILE\_WIDTH times

- Given 48KB shared memory:
  - Max tiled size =  $(48\text{KB}/4\text{B}/3)^{(1/2)} = 64$
  - CGMA = number of data re-use = TILE\_WIDTH = 64!

```

extern __shared__ int S[];
inline int Addr(int matrixIdx, int i, int j, int N) {
    return (N*N*matrixIdx + i*N+ j);
}
_global_ void Matrix_Mul(int* A, int* B,int* C, int* N) {
    int i = threadIdx.x;
    int j = threadIdx.y;
    //move data to shared memory
    S[Addr(0, i, j, N)]=A[Addr(0, i, j, N)];
    S[Addr(1, i, j, N)]=B[Addr(0, i, j, N)];
    __syncthreads();
    // do computation
    for(int k=0; k<*N; k++)
        C[Addr(1, i, j, N)]=S[Addr(0, i, k, N)]*S[Addr(1, k, j, N)];
}
void main() {
    for(int i=0; i<N; i+=TILE_WIDTH)
        for(int j=0; j<N; j+=TILE_WIDTH){
            cudaMemcpy(d_A, &A[i,j], sizeof(int)*TILE_WIDTH*TILE_WIDTH, H2D);
            cudaMemcpy(d_B, &B[i,j], sizeof(int)*TILE_WIDTH*TILE_WIDTH, H2D);
            Matrix_Mul<<<1, N, 2*N*N>>>(d_A, d_B, d_C, TILE_WIDTH);
            cudaMemcpy(&C[i,j], d_C, sizeof(int)*TILE_WIDTH*TILE_WIDTH), D2H;
        }
}

```

# Tiled Algorithm

- Block algorithms or tiled algorithms:
  - Split the inputs into blocks to fit into shared (cache) memory
  - Increase data reuse, minimize global memory access
- Larger CGMA ratio does not always guarantee better performance.
  - CGMA ratio should be large enough to hide the communication cost, not the larger the better
  - Block algorithms cause overhead due to increasing computations or number of thread blocks

# Outline

- Host memory
  - Pined memory
  - Asynchronous computation & data transfer
  - Streams
- Global/Local memory
  - Memory coalescing
  - Tiled algorithm
- Shared memory
  - Bank conflicts avoidance
  - Memory padding
- Address linearization

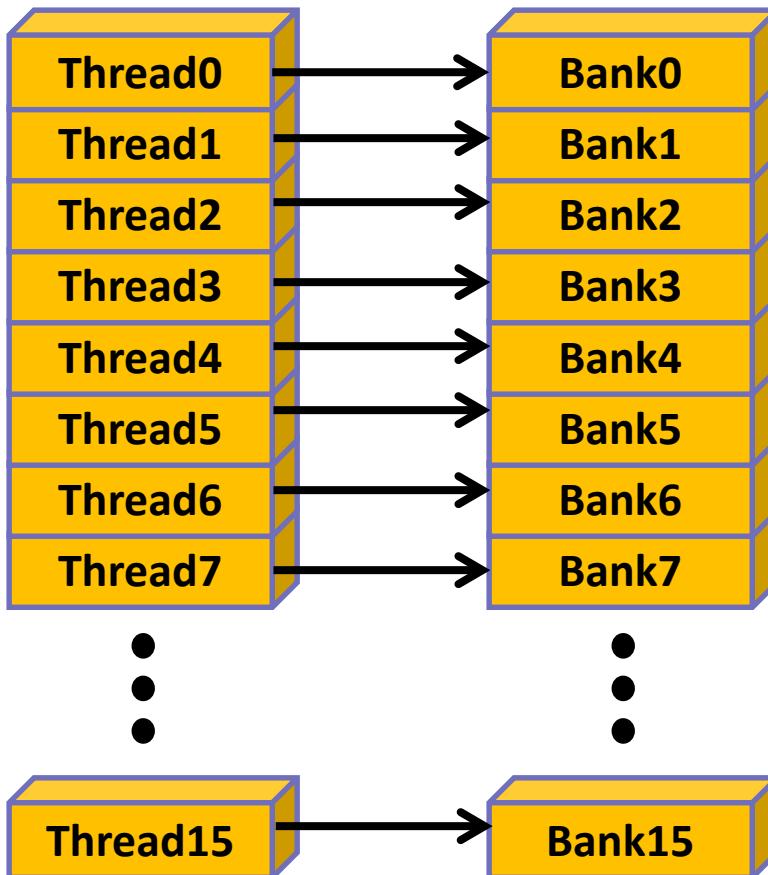
# Shared Memory Architecture

- Many threads accessing memory
  - Therefore, memory is divided into **banks**
  - Successive 32-bit (4Bytes) words assigned to successive banks
- Each bank can service one address per cycle
  - A memory can service as many simultaneous accesses as it has banks
- Multiple simultaneous accesses to a bank result in a **bank conflict**
  - Conflicting accesses are serialized
- Shared memory is as fast as register if no bank conflict

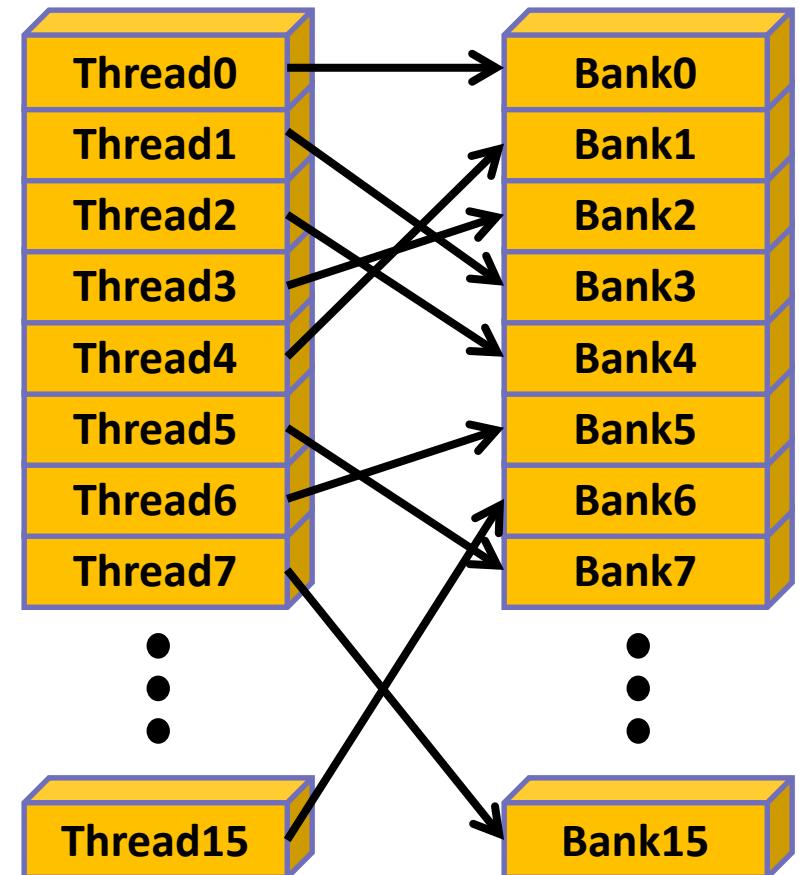


# Example: No bank Conflict

## ■ Linear addressing

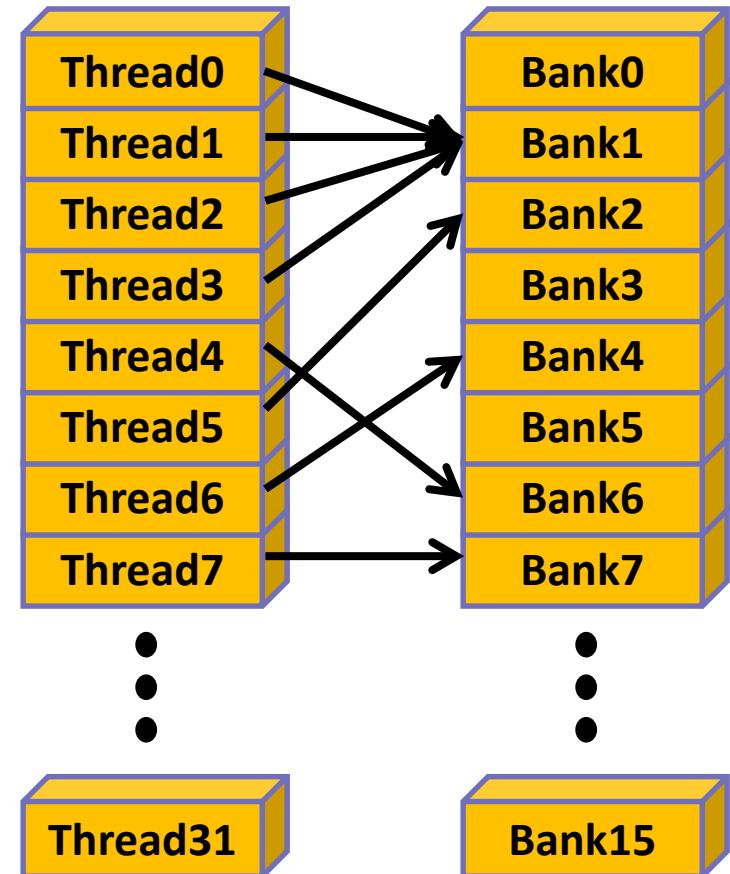


## ■ Random 1:1 Permutation



# Example: No bank Conflict

- If all threads of a **half-warp** read the identical address, there is no bank conflict (using broadcast)
  - Assume warp size is 8
  - Thread0~3 access the same data & in the same half-warp
  - The rest of threads also have 1:1 permutation and no conflict
  - But not for write access



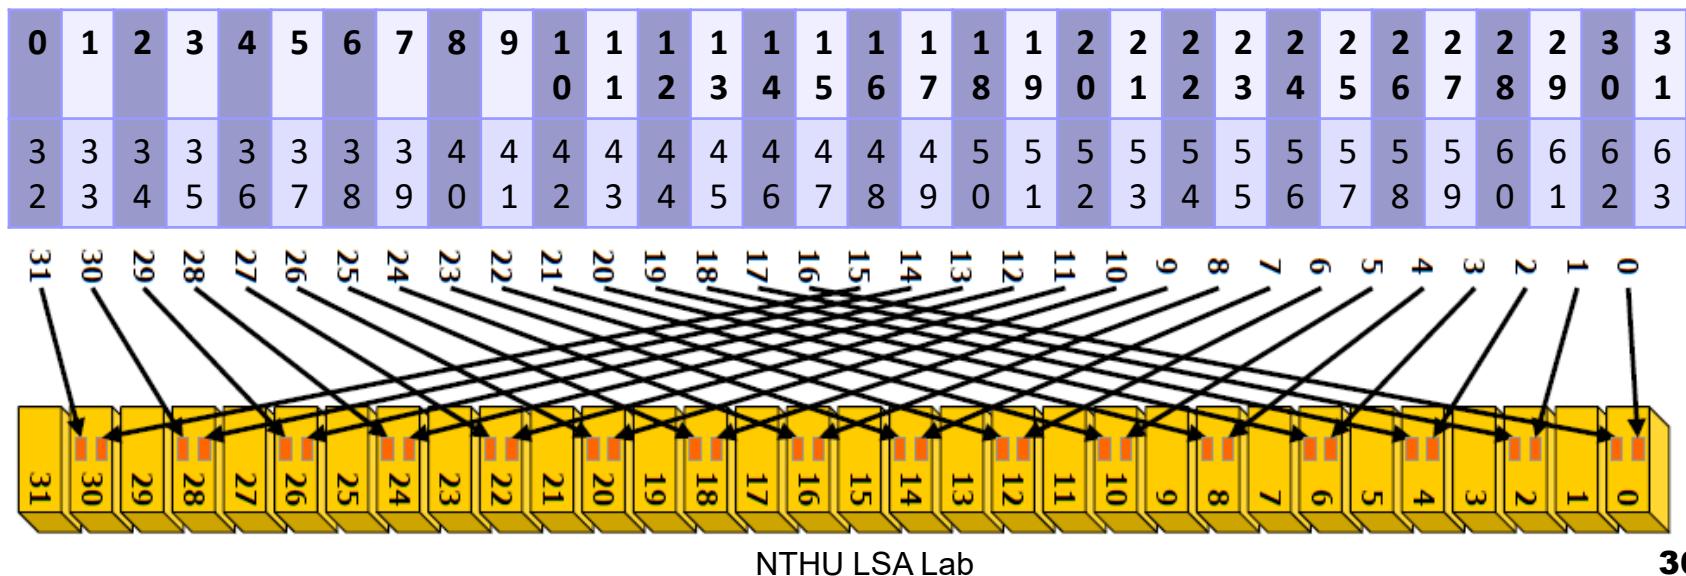
# Example: Bank Conflict

## ■ n-way bank conflict

- Each bank has n different memory access

## ■ Ex: 2-way bank conflict

```
__shared__ int array[2][32];
int offset = threadIdx.x*2;
int temp = array[offset/32][offset%32];
```



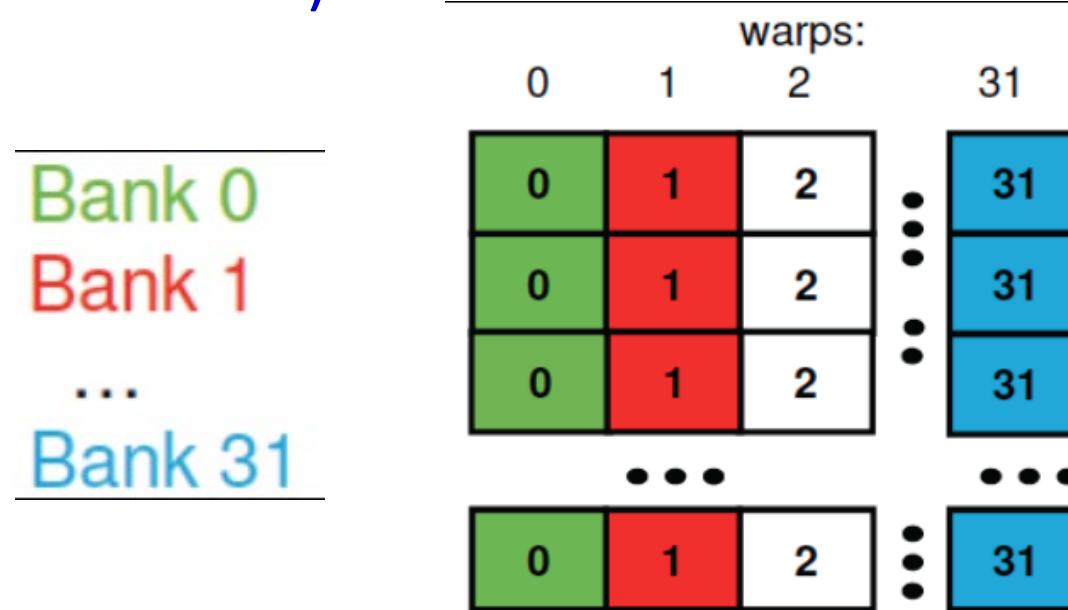
# Bank Conflict Avoidance

- Change shared memory access pattern
  - Linear addressing access
  - 1:1 permutation
  - Broadcast: half-warp read the identical address
- Memory padding
  - Add addition memory space to avoid bank conflict

# Example: 2D array

## ■ 32x32 SMEM array

- Warp accesses a column:
- 32-way bank conflicts (threads in a warp access the same bank)



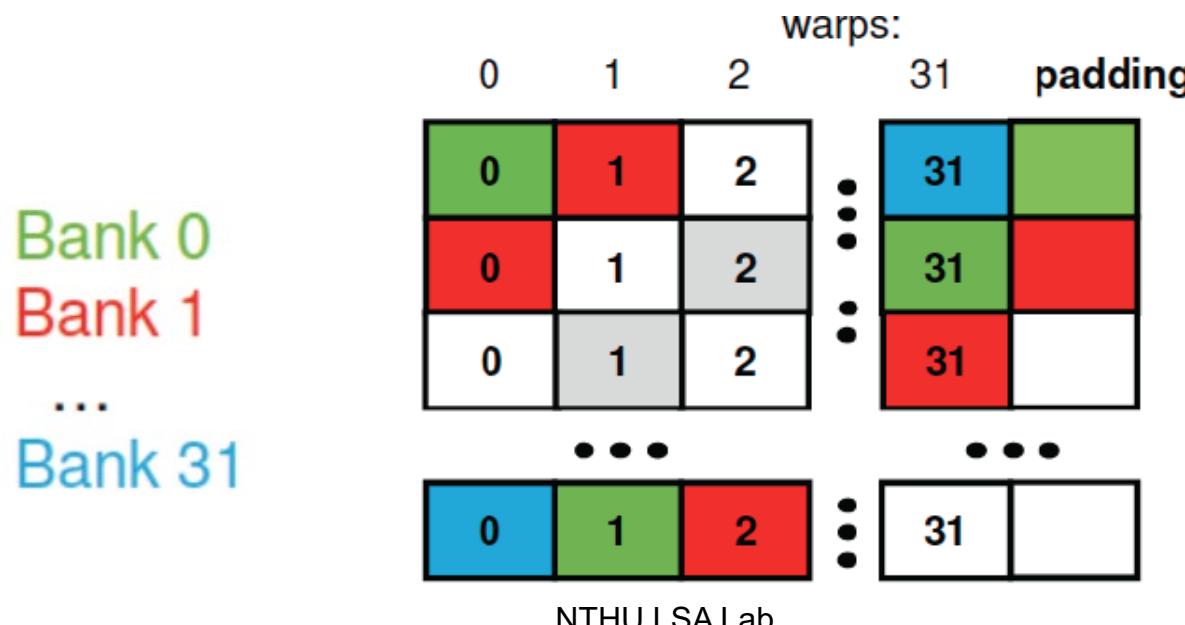
# Memory Padding

- Add a column for padding:

- 32x33 SMEM array

- Warp accesses a column:

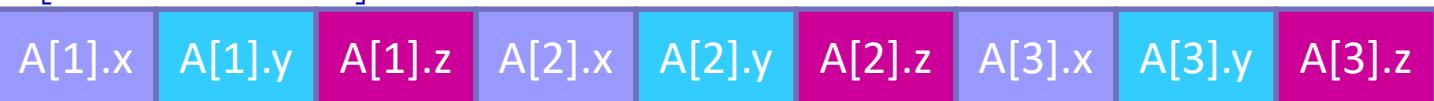
- 32 different banks, no bank conflicts



# Address linearization (SoA)

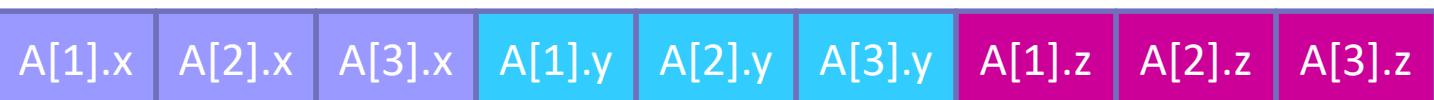
- Address linearization can avoid bank conflict on shared memory, and provide memory coalescing on local memory or constant memory
- An array of structures behaves like row major accesses

- struct Point { double x; double y; double z; } A[N];
- A[threadIdx].x = ...



- A structure of arrays behaves like column major

- struct PointList{double \*x; double \*y; double \*z;} A;
- A.x[threadIdx] = ...





Slides from Mark Harris, NVIDIA Developer Technology  
Performance Optimization

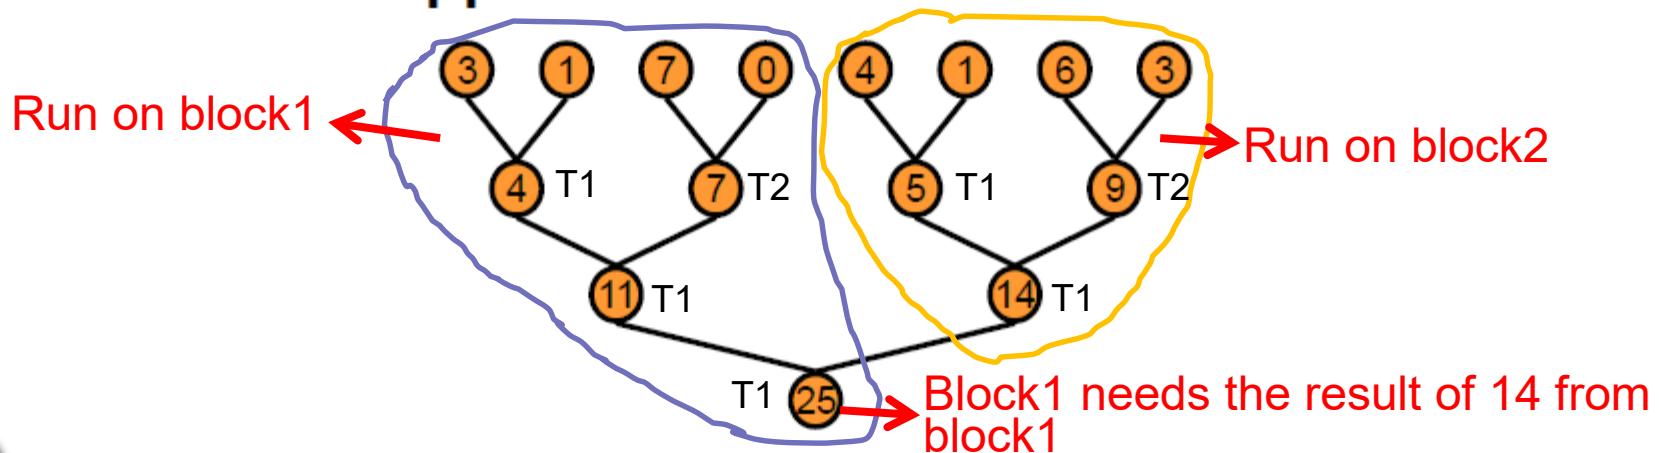
# **AN EXAMPLE OF CUDA**

# Parallel Reduction

- ➊ Common and important data parallel primitive
- ➋ Easy to implement in CUDA
  - ➌ Harder to get it right      Performance!
- ➌ Serves as a great optimization example      30x Speedup!
  - ➍ We'll walk step by step through 7 different versions
  - ➎ Demonstrates several important optimization strategies

# Parallel Reduction

- Tree-based approach used within each thread block



- Need to be able to use multiple thread blocks
  - To process very large arrays
  - To keep all multiprocessors on the GPU busy
  - Each thread block reduces a portion of the array
- But how do we communicate partial results between thread blocks?

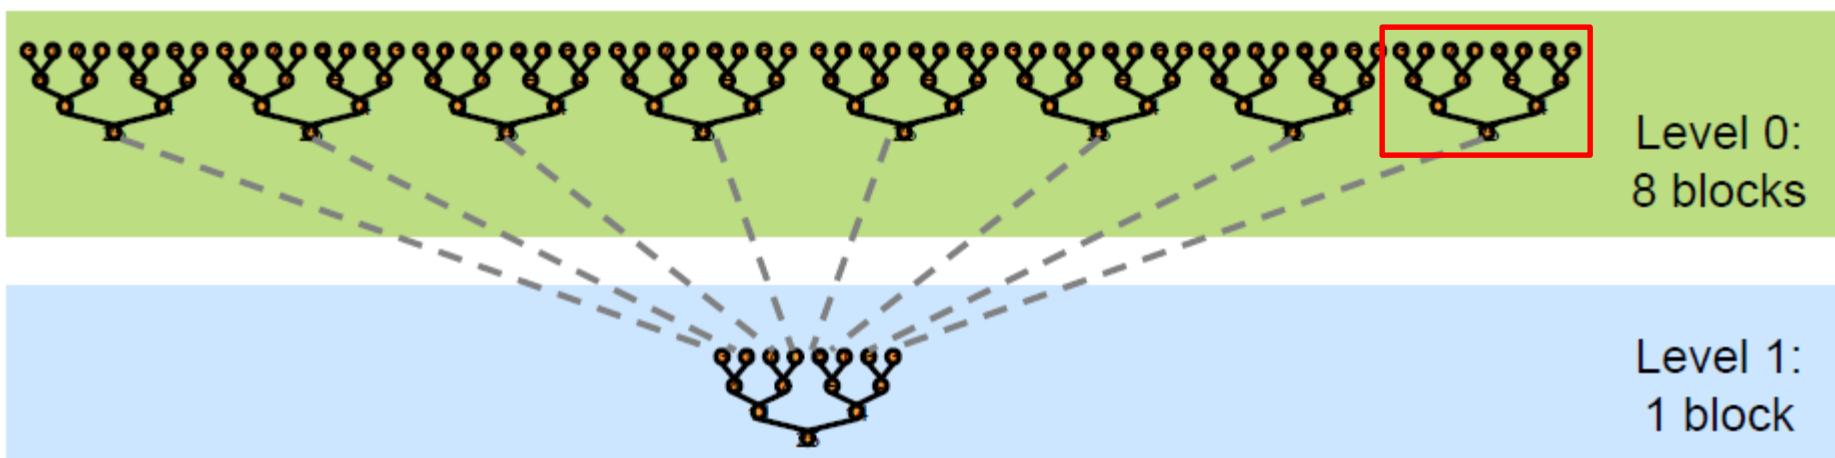
# Problem: Global Synchronization

- If we could synchronize across all thread blocks, could easily reduce very large arrays, right?
  - Global sync after each block produces its result
  - Once all blocks reach sync, continue recursively
- But CUDA has no global synchronization. Why?
  - Expensive to build in hardware for GPUs with high processor count
  - Would force programmer to run fewer blocks (no more than  $\# \text{ multiprocessors} * \# \text{ resident blocks} / \text{multiprocessor}$ ) to avoid deadlock, which may reduce overall efficiency
- Solution: decompose into multiple kernels
  - Kernel launch serves as a global synchronization point
  - Kernel launch has negligible HW overhead, low SW overhead

# Solution: Kernel Decomposition

- Avoid global sync by decomposing computation into multiple kernel invocations

If the maximum threads per block is 8



- In the case of reductions, code for all levels is the same
  - Recursive kernel invocation



# Reduction #1: Interleaved Addressing

// input/output data is initiated on global memory

```
__global__ void reduce0(int *g_idata, int *g_odata) {  
    extern __shared__ int sdata[]; // Use shared memory for computations
```

// each thread loads one element from global to shared mem

```
    unsigned int tid = threadIdx.x;  
    unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;  
    sdata[tid] = g_idata[i];  
    __syncthreads(); // Wait for other threads to finish moving
```

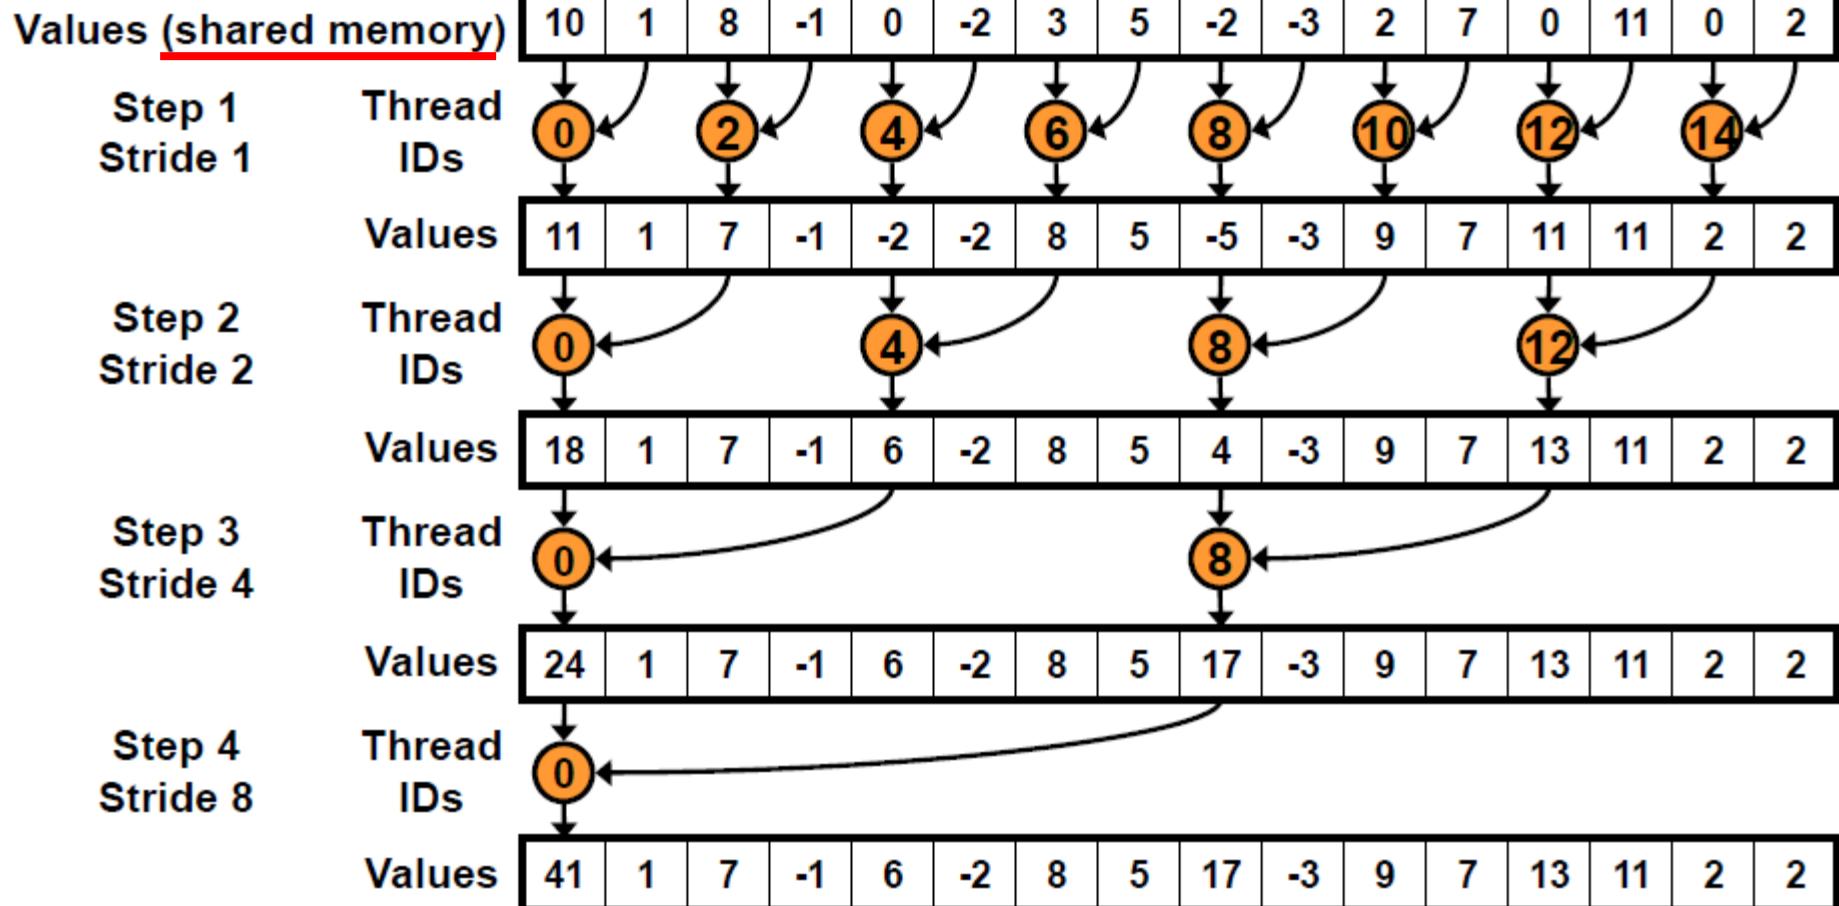
// do reduction in shared mem

```
    for(unsigned int s=1; s < blockDim.x; s *= 2) {  
        if (tid % (2*s) == 0) {  
            sdata[tid] += sdata[tid + s];  
        }  
        __syncthreads(); // Sync between threads in the same block  
    }
```

// write result for this block to global mem

```
    if (tid == 0) g_odata[blockIdx.x] = sdata[0];  
}
```

# Parallel Reduction: Interleaved Addressing



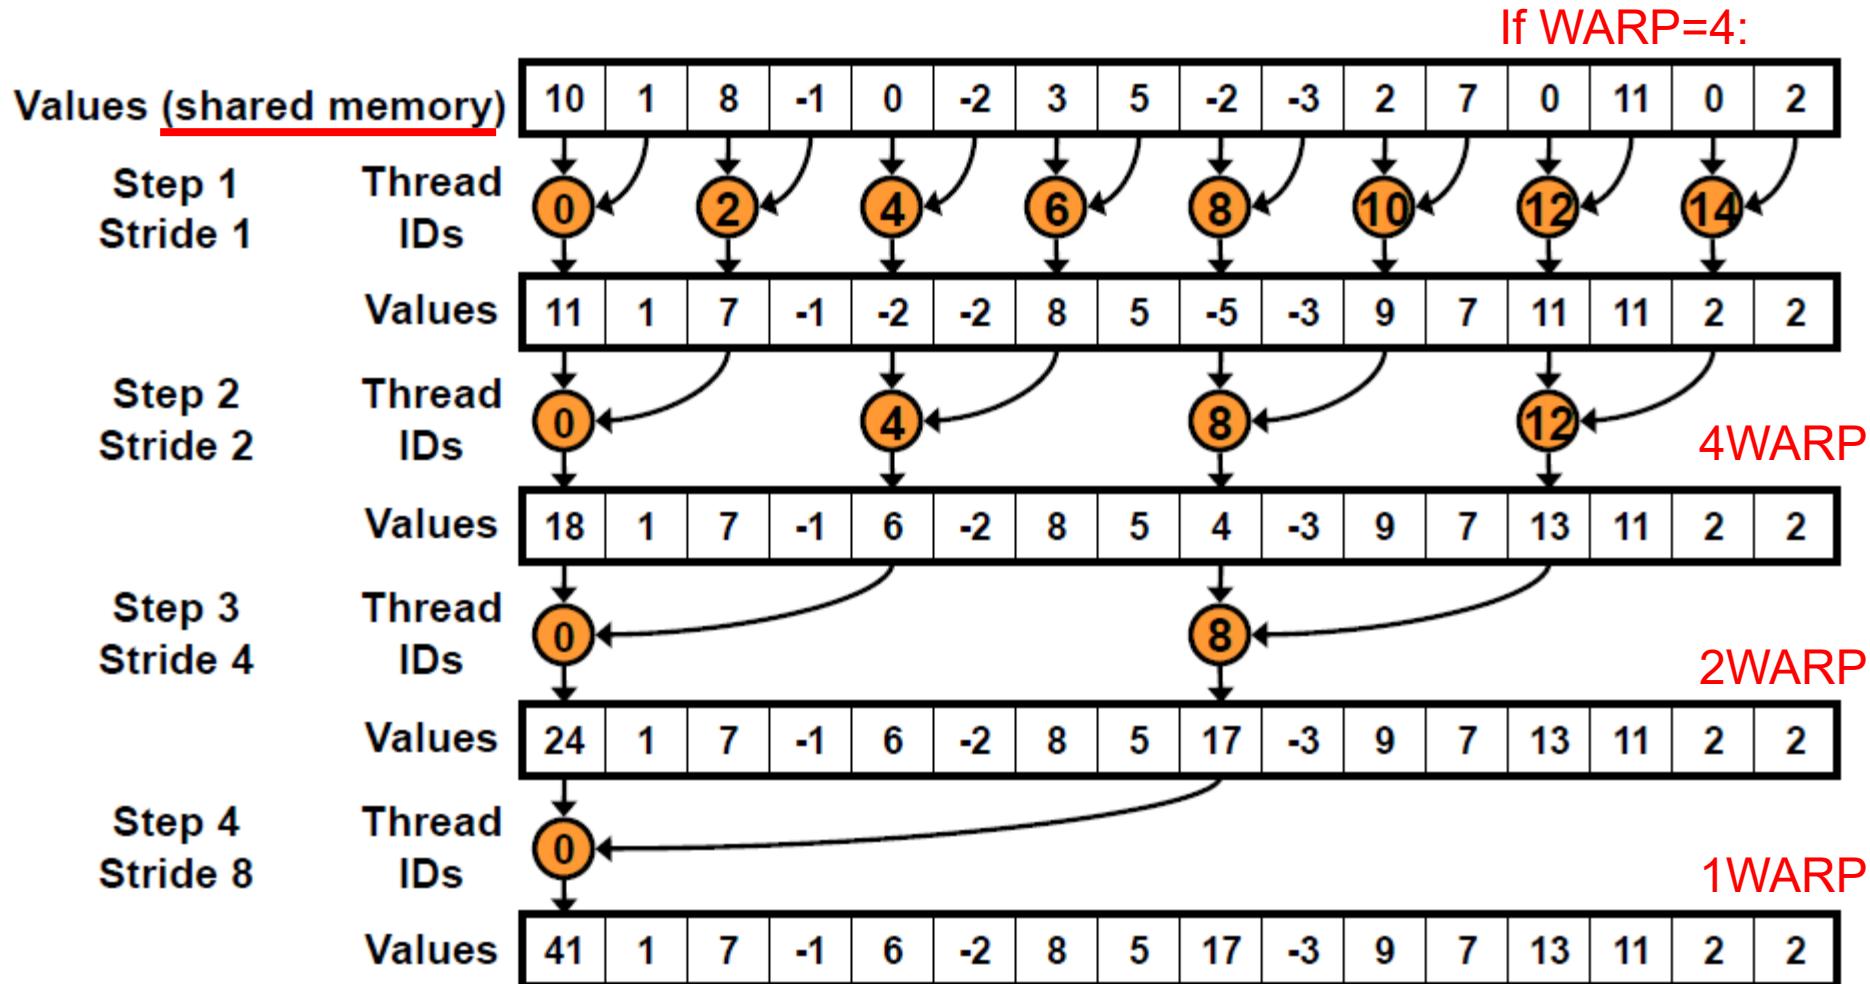
# Performance for 4M element reduction



	Time ( $2^{22}$ ints)	Bandwidth
<b>Kernel 1:</b> interleaved addressing with divergent branching	<b>8.054 ms</b>	<b>2.083 GB/s</b>

Note: Block Size = 128 threads for all tests

# Parallel Reduction: Interleaved Addressing



Highly divergent wrap (threadID 0~14)

NTHU LSA Lab

# Reduction #1: Interleaved Addressing



```
__global__ void reduce1(int *g_idata, int *g_odata) {
    extern __shared__ int sdata[];

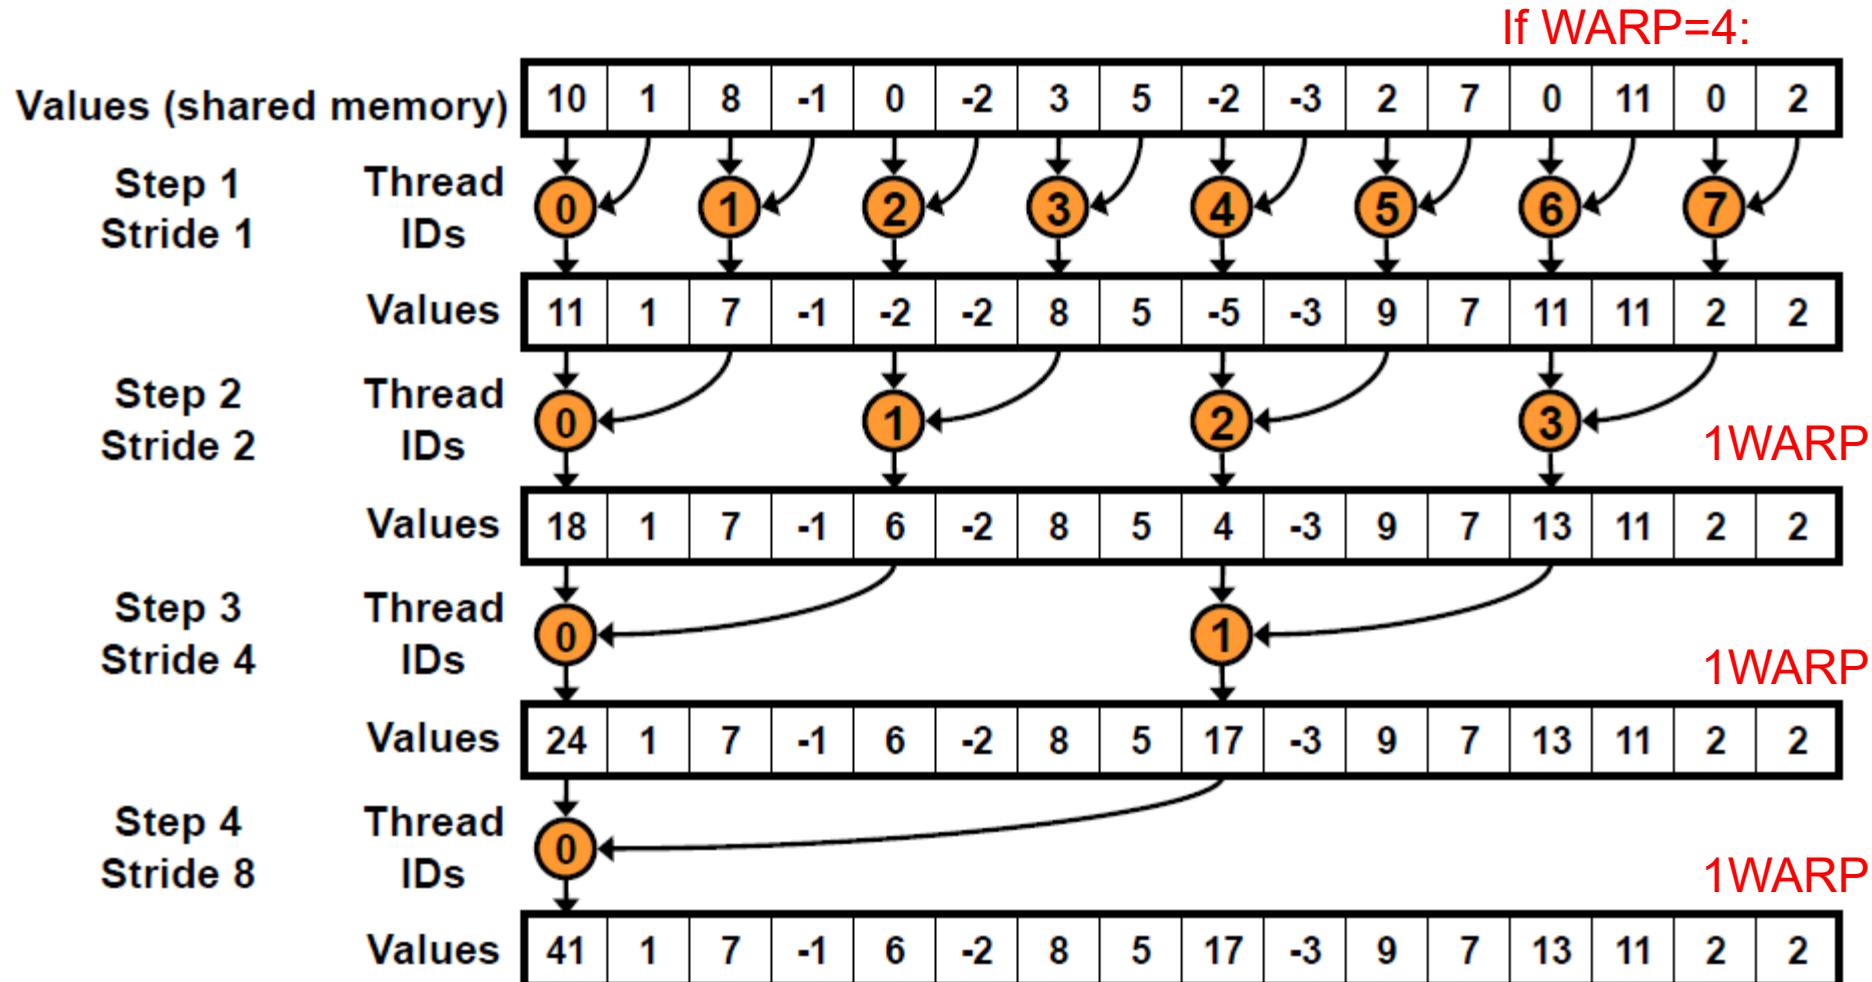
    // each thread loads one element from global to shared mem
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
    sdata[tid] = g_idata[i];
    __syncthreads();
```

```
// do reduction in shared mem
for (unsigned int s=1; s < blockDim.x; s *= 2) {
    if (tid % (2*s) == 0) { ←
        sdata[tid] += sdata[tid + s];
    }
    __syncthreads();
}
```

Problem: highly divergent warps are very inefficient, and % operator is very slow

```
// write result for this block to global mem
if (tid == 0) g_odata[blockIdx.x] = sdata[0];
}
```

# Parallel Reduction: Interleaved Addressing



# Reduction #2: Interleaved Addressing

Just replace divergent branch in inner loop:

```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    if (tid % (2*s) == 0) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```

With strided index and non-divergent branch:

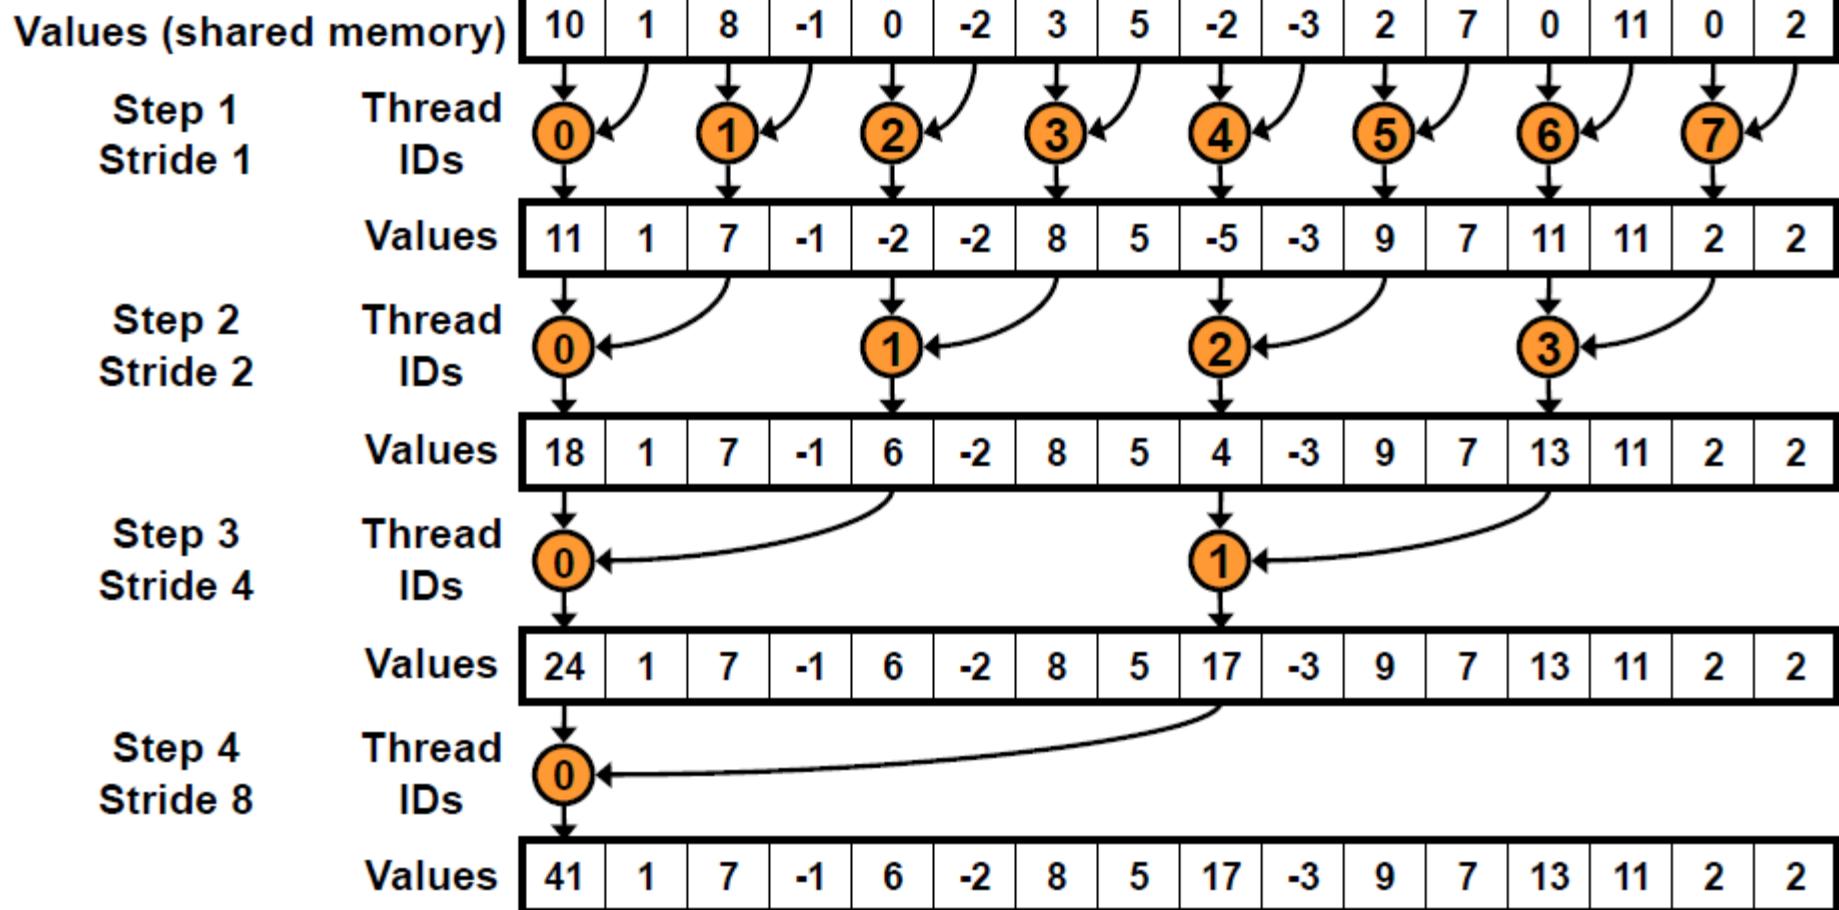
```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    int index = 2 * s * tid;  
  
    if (index < blockDim.x) {  
        sdata[index] += sdata[index + s];  
    }  
    __syncthreads();  
}
```

# Performance for 4M element reduction



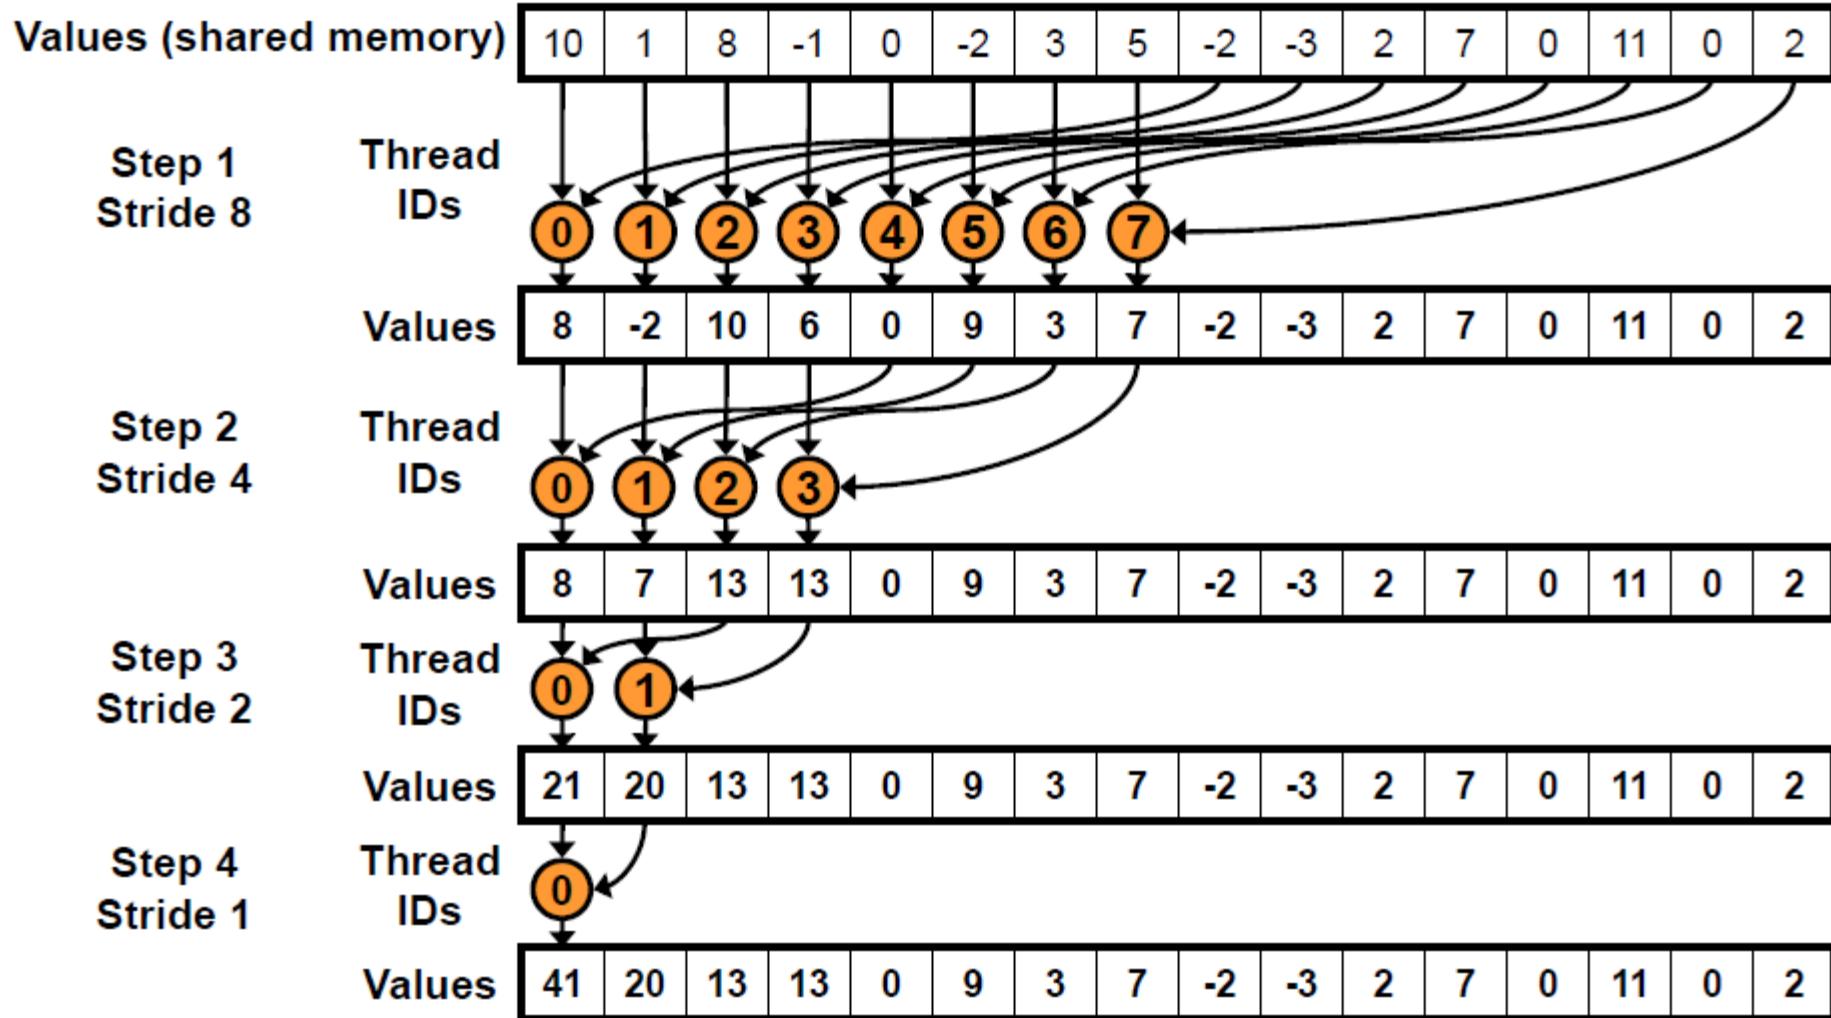
	Time ( $2^{22}$ ints)	Bandwidth	Step Speedup	Cumulative Speedup
<b>Kernel 1:</b> interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
<b>Kernel 2:</b> interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x

# Parallel Reduction: Interleaved Addressing



New Problem: Shared Memory Bank Conflicts

# Parallel Reduction: Sequential Addressing



Sequential addressing is conflict free

# Reduction #3: Sequential Addressing

Just replace strided indexing in inner loop:

```
for (unsigned int s=1; s < blockDim.x; s *= 2) {  
    int index = 2 * s * tid;  
  
    if (index < blockDim.x) {  
        sdata[index] += sdata[index + s];  
    }  
    __syncthreads();  
}
```

With reversed loop and threadID-based indexing:

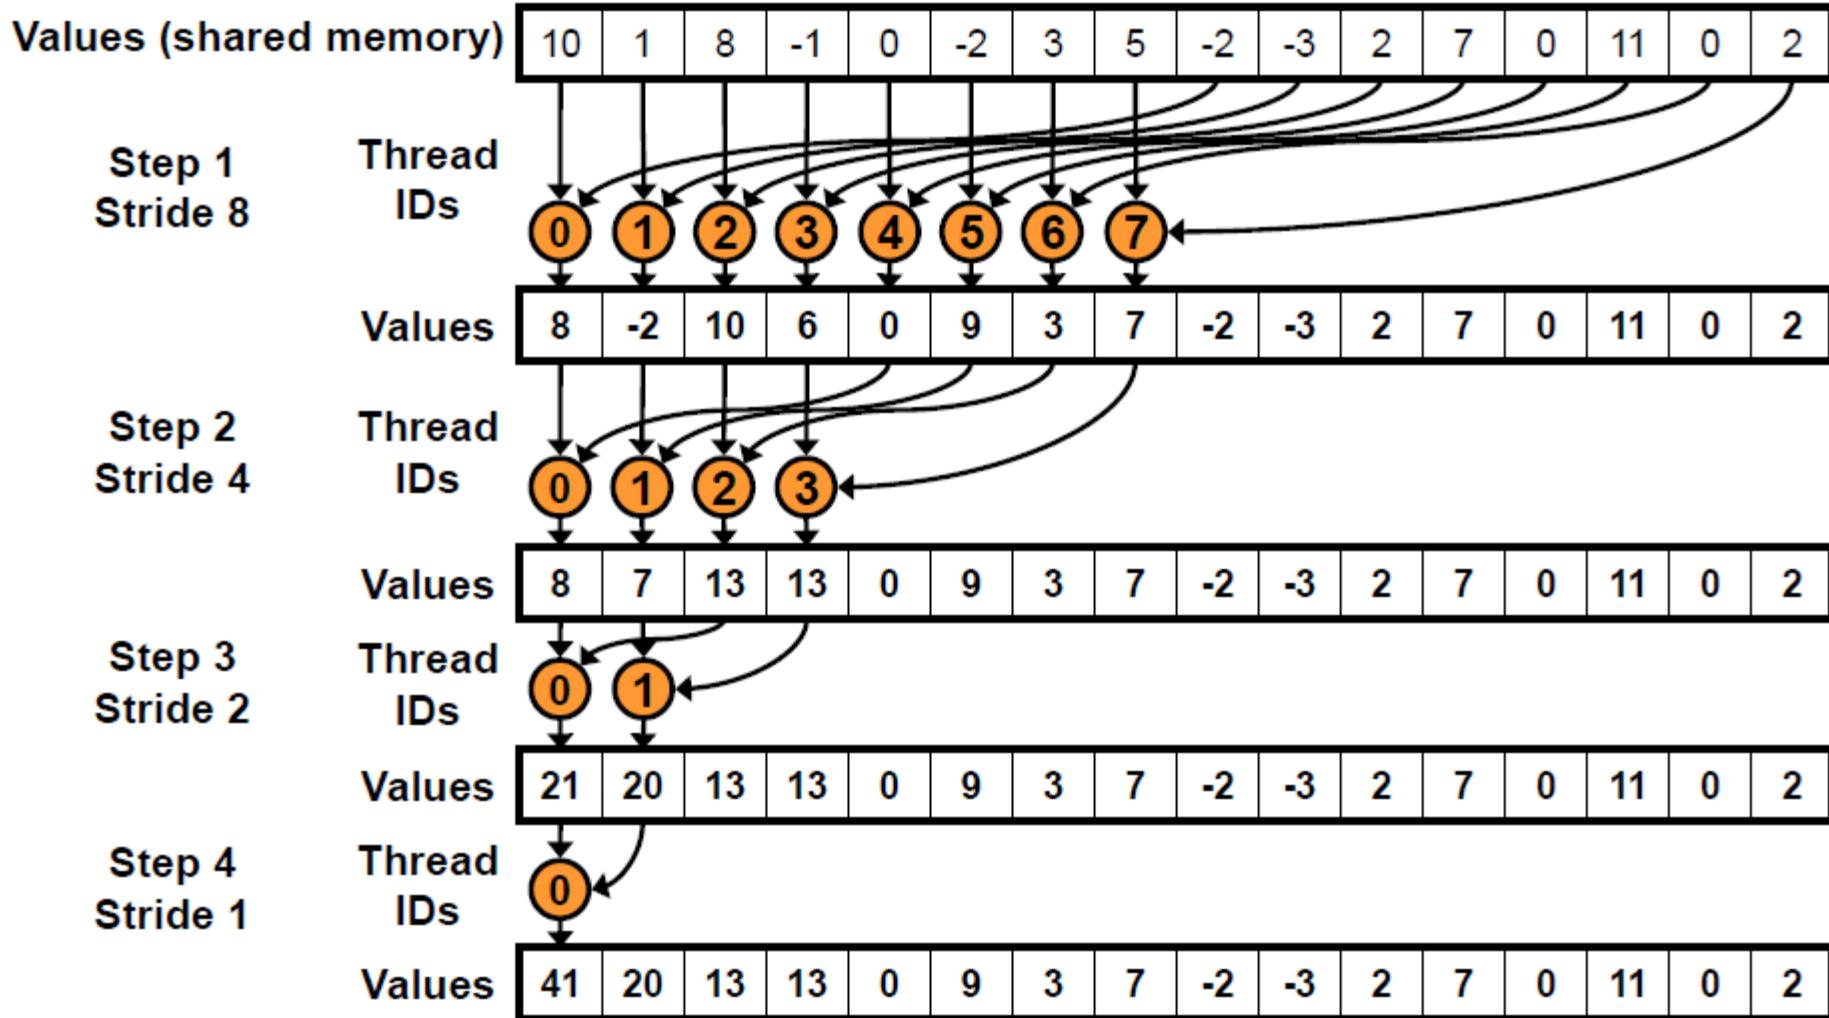
```
for (unsigned int s=blockDim.x/2; s>0; s>>=1) {  
    if (tid < s) {  
        sdata[tid] += sdata[tid + s];  
    }  
    __syncthreads();  
}
```

# Performance for 4M element reduction



	Time ( $2^{22}$ ints)	Bandwidth	Step Speedup	Cumulative Speedup
<b>Kernel 1:</b> interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
<b>Kernel 2:</b> interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
<b>Kernel 3:</b> sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x

# Parallel Reduction: Sequential Addressing



Half of the threads are idle since 1<sup>st</sup> iteration!



# Reduction #4: First Add During Load

Halve the number of blocks, and replace single load:

```
// each thread loads one element from global to shared mem
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
sdata[tid] = g_idata[i];
__syncthreads();
```

With two loads and first add of the reduction:

```
// perform first level of reduction,
// reading from global memory, writing to shared memory
unsigned int tid = threadIdx.x;
unsigned int i = blockIdx.x*(blockDim.x*2) + threadIdx.x;
sdata[tid] = g_idata[i] + g_idata[i+blockDim.x];
__syncthreads();
```

# Performance for 4M element reduction



	Time ( $2^{22}$ ints)	Bandwidth	Step Speedup	Cumulative Speedup
<b>Kernel 1:</b> interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
<b>Kernel 2:</b> interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
<b>Kernel 3:</b> sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
<b>Kernel 4:</b> first add during global load	0.965 ms	17.377 GB/s	1.78x	8.34x

# Performance for 4M element reduction



	Time ( $2^{22}$ ints)	Bandwidth	Step Speedup	Cumulative Speedup
<b>Kernel 1:</b> interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
<b>Kernel 2:</b> interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
<b>Kernel 3:</b> sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
<b>Kernel 4:</b> first add during global load	0.965 ms	17.377 GB/s	1.78x	8.34x
<b>Kernel 5:</b> unroll last warp	0.536 ms	31.289 GB/s	1.8x	15.01x
<b>Kernel 6:</b> completely unrolled	0.381 ms	43.996 GB/s	1.41x	21.16x
<b>Kernel 7:</b> multiple elements per thread	0.268 ms	62.671 GB/s	1.42x	30.04x

Details in backup slides

Kernel 7 on 32M elements: 73 GB/s!  
NTHU LSA Lab



```
template <unsigned int blockSize>
__device__ void warpReduce(volatile int *sdata, unsigned int tid) {
    if (blockSize >= 64) sdata[tid] += sdata[tid + 32];
    if (blockSize >= 32) sdata[tid] += sdata[tid + 16];
    if (blockSize >= 16) sdata[tid] += sdata[tid + 8];
    if (blockSize >= 8) sdata[tid] += sdata[tid + 4];
    if (blockSize >= 4) sdata[tid] += sdata[tid + 2];
    if (blockSize >= 2) sdata[tid] += sdata[tid + 1];
}

template <unsigned int blockSize>
__global__ void reduce6(int *g_idata, int *g_odata, unsigned int n) {
    extern __shared__ int sdata[];
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x*(blockSize*2) + tid;
    unsigned int gridSize = blockSize*2*gridDim.x;
    sdata[tid] = 0;

    while (i < n) { sdata[tid] += g_idata[i] + g_idata[i+blockSize]; i += gridSize; }
    __syncthreads();

    if (blockSize >= 512) { if (tid < 256) { sdata[tid] += sdata[tid + 256]; } __syncthreads(); }
    if (blockSize >= 256) { if (tid < 128) { sdata[tid] += sdata[tid + 128]; } __syncthreads(); }
    if (blockSize >= 128) { if (tid < 64) { sdata[tid] += sdata[tid + 64]; } __syncthreads(); }

    if (tid < 32) warpReduce(sdata, tid);
    if (tid == 0) g_odata[blockIdx.x] = sdata[0];
}
```

## Final Optimized Kernel

# Backup

# Instruction Bottleneck

- ➊ At 17 GB/s, we're far from bandwidth bound
  - ➌ And we know reduction has low arithmetic intensity
- ➋ Therefore a likely bottleneck is instruction overhead
  - ➌ Ancillary instructions that are not loads, stores, or arithmetic for the core computation
  - ➌ In other words: address arithmetic and loop overhead
- ➌ Strategy: unroll loops

# Unrolling the Last Warp

- ➊ As reduction proceeds, # “active” threads decreases
  - ➌ When  $s \leq 32$ , we have only one warp left
- ➋ Instructions are SIMD synchronous within a warp
- ➌ That means when  $s \leq 32$ :
  - ➌ We don’t need to `_syncthreads()`
  - ➌ We don’t need “if ( $tid < s$ )” because it doesn’t save any work
- ➍ Let’s unroll the last 6 iterations of the inner loop

# Reduction #5: Unroll the Last Warp

```
__device__ void warpReduce(volatile int* sdata, int tid) {  
    sdata[tid] += sdata[tid + 32];  
    sdata[tid] += sdata[tid + 16];  
    sdata[tid] += sdata[tid + 8];  
    sdata[tid] += sdata[tid + 4];  
    sdata[tid] += sdata[tid + 2];  
    sdata[tid] += sdata[tid + 1];  
}
```

IMPORTANT:  
For this to be correct,  
we must use the  
“volatile” keyword!

```
// later...  
for (unsigned int s=blockDim.x/2; s>32; s>>=1) {  
    if (tid < s)  
        sdata[tid] += sdata[tid + s];  
    __syncthreads();  
}  
  
if (tid < 32) warpReduce(sdata, tid);
```

**Note: This saves useless work in *all* warps, not just the last one!**

Without unrolling, all warps execute every iteration of the for loop and if statement

# Performance for 4M element reduction



	Time ( $2^{22}$ ints)	Bandwidth	Step Speedup	Cumulative Speedup
<b>Kernel 1:</b> interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
<b>Kernel 2:</b> interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
<b>Kernel 3:</b> sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
<b>Kernel 4:</b> first add during global load	0.965 ms	17.377 GB/s	1.78x	8.34x
<b>Kernel 5:</b> unroll last warp	0.536 ms	31.289 GB/s	1.8x	15.01x

# Complete Unrolling

- ➊ If we knew the number of iterations at compile time, we could completely unroll the reduction
  - ➌ Luckily, the block size is limited by the GPU to 512 threads
  - ➌ Also, we are sticking to power-of-2 block sizes
- ➋ So we can easily unroll for a fixed block size
  - ➌ But we need to be generic – how can we unroll for block sizes that we don't know at compile time?
- ➌ Templates to the rescue!
  - ➌ CUDA supports C++ template parameters on device and host functions

# Unrolling with Templates

- Specify block size as a function template parameter:

```
template <unsigned int blockSize>
__global__ void reduce5(int *g_idata, int *g_odata)
```

# Reduction #6: Completely Unrolled

```
Template <unsigned int blockSize>
__device__ void warpReduce(volatile int* sdata, int tid) {
    if (blockSize >= 64) sdata[tid] += sdata[tid + 32];
    if (blockSize >= 32) sdata[tid] += sdata[tid + 16];
    if (blockSize >= 16) sdata[tid] += sdata[tid + 8];
    if (blockSize >= 8) sdata[tid] += sdata[tid + 4];
    if (blockSize >= 4) sdata[tid] += sdata[tid + 2];
    if (blockSize >= 2) sdata[tid] += sdata[tid + 1];
}
```

```
if (blockSize >= 512) {
    if (tid < 256) { sdata[tid] += sdata[tid + 256]; } __syncthreads();
}
if (blockSize >= 256) {
    if (tid < 128) { sdata[tid] += sdata[tid + 128]; } __syncthreads();
}
if (blockSize >= 128) {
    if (tid < 64) { sdata[tid] += sdata[tid + 64]; } __syncthreads();
}

if (tid < 32) warpReduce<blockSize>(sdata, tid);
```

Note: all code in RED will be evaluated at compile time.

Results in a very efficient inner loop!

# Invoking Template Kernels

- Don't we still need block size at compile time?

- Nope, just a switch statement for 10 possible block sizes:

```
switch (threads)
{
    case 512:
        reduce5<512><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 256:
        reduce5<256><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 128:
        reduce5<128><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 64:
        reduce5< 64><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 32:
        reduce5< 32><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 16:
        reduce5< 16><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 8:
        reduce5< 8><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 4:
        reduce5< 4><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 2:
        reduce5< 2><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
    case 1:
        reduce5< 1><<< dimGrid, dimBlock, smemSize >>>(d_idata, d_odata); break;
}
```

# Performance for 4M element reduction



	Time ( $2^{22}$ ints)	Bandwidth	Step Speedup	Cumulative Speedup
<b>Kernel 1:</b> interleaved addressing with divergent branching	8.054 ms	2.083 GB/s		
<b>Kernel 2:</b> interleaved addressing with bank conflicts	3.456 ms	4.854 GB/s	2.33x	2.33x
<b>Kernel 3:</b> sequential addressing	1.722 ms	9.741 GB/s	2.01x	4.68x
<b>Kernel 4:</b> first add during global load	0.965 ms	17.377 GB/s	1.78x	8.34x
<b>Kernel 5:</b> unroll last warp	0.536 ms	31.289 GB/s	1.8x	15.01x
<b>Kernel 6:</b> completely unrolled	0.381 ms	43.996 GB/s	1.41x	21.16x

# Reference

- NVIDIA Advanced CUDA Webinar Memory Optimizations
  - [http://on-demand.gputechconf.com/gtc-express/2011/presentations/NVIDIA\\_GPU\\_Computing\\_Webinars\\_CUDA\\_Memory\\_Optimization.pdf](http://on-demand.gputechconf.com/gtc-express/2011/presentations/NVIDIA_GPU_Computing_Webinars_CUDA_Memory_Optimization.pdf)
- NVIDIA CUDA C/C++ Streams and Concurrency
  - <http://on-demand.gputechconf.com/gtc-express/2011/presentations/StreamsAndConcurrencyWebinar.pdf>
- Mark Harris, NVIDIA Developer Technology
  - [http://gpgpu.org/static/sc2007/SC07\\_CUDA\\_5\\_Optimization\\_Harris.pdf](http://gpgpu.org/static/sc2007/SC07_CUDA_5_Optimization_Harris.pdf)