

Lab3 CUDA Basic

Oct, 2025 Parallel Programming

Overview

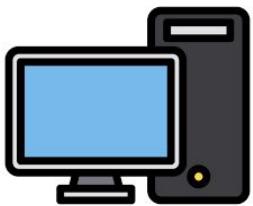
- ❖ Platform guide
- ❖ Tools
- ❖ Assignment

Platform Guide

The GPU Cluster

- ❖ Host: **apollo.cs.nthu.edu.tw**
- ❖ Account & Password: Same as apollo origo

Client
(Your computer)

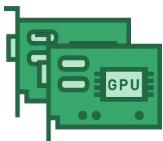


`ssh apollo.cs.nthu.edu.tw`

Frontend Server
without GPU
(**apollo-login**)

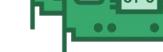
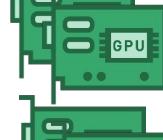
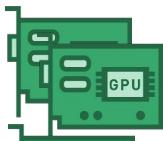


`ssh nv-test`

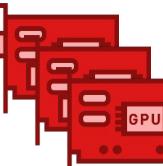
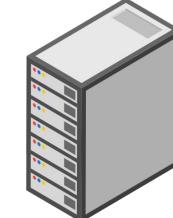


`srun -p nvidia`

`nv-vm[1-7]`

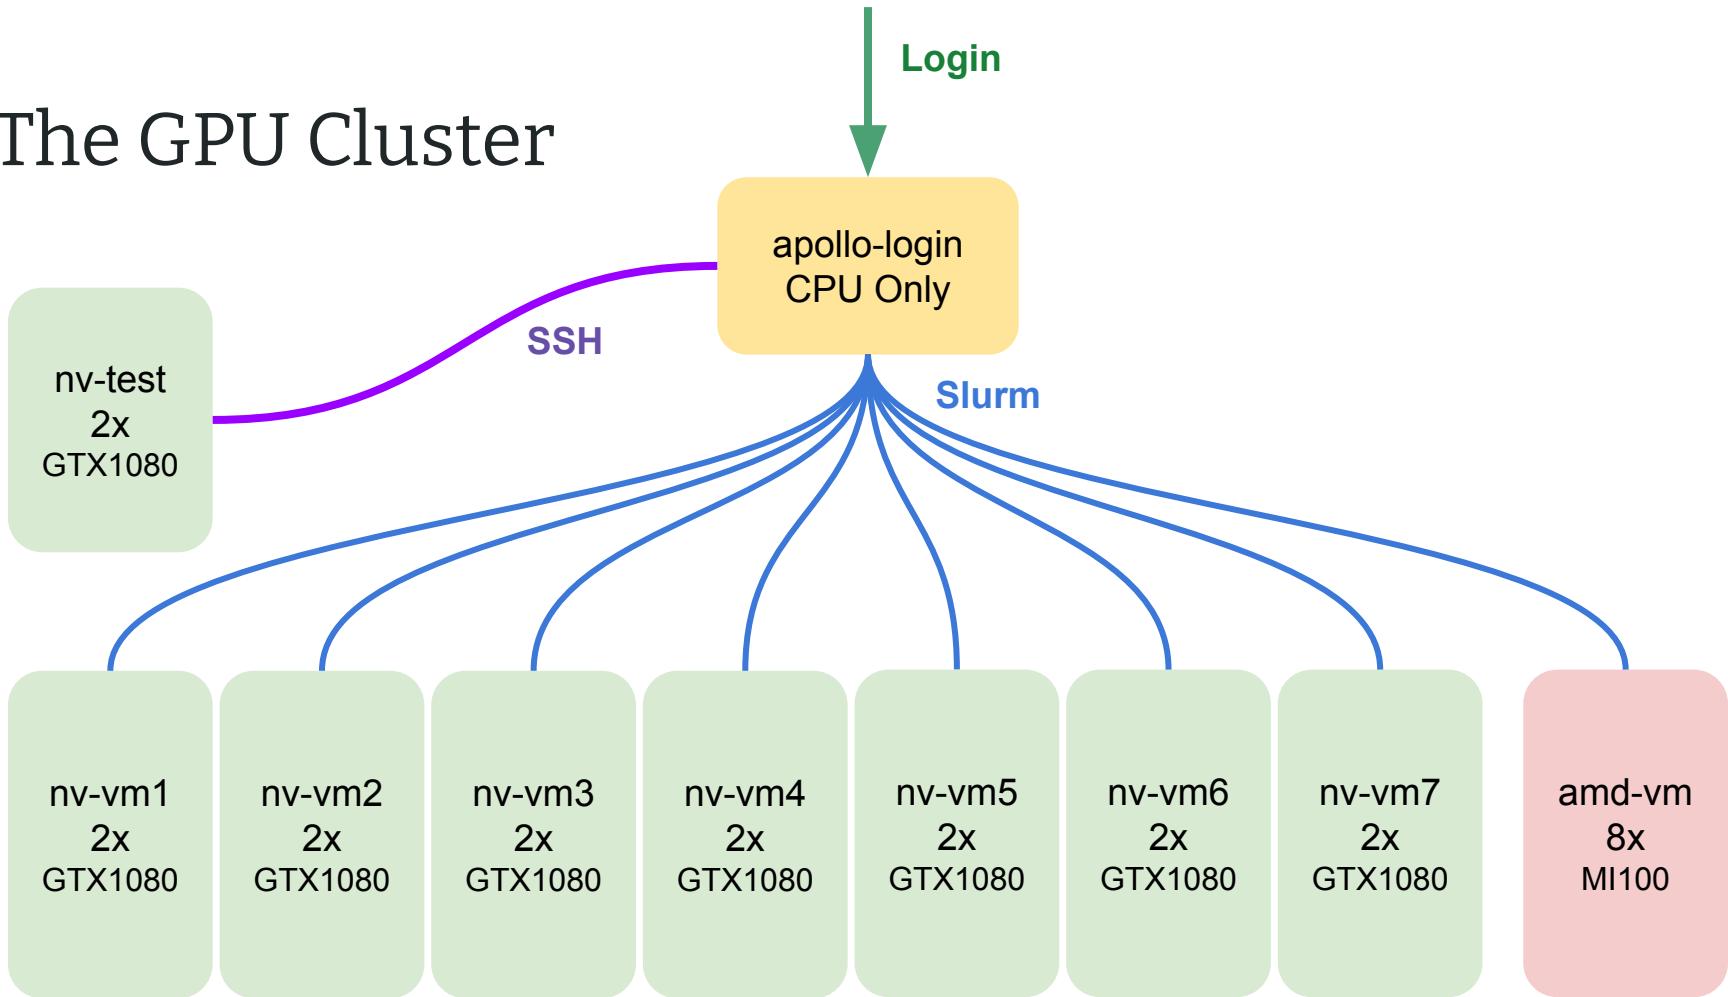


`srun -p amd`
`amd-vm`



8 NVIDIA GPU VMs: each with 2x GTX 1080 GPUs
2 AMD GPU VM: with 8x AMD Instinct MI100 GPUs

The GPU Cluster



Job Scheduler

- ❖ Slurm
- ❖ Partitions: **nvidia** for NVIDIA GPUs (default), **amd** for AMD GPUs

```
gilbert12@apollo-login:~$ sinfo
PARTITION    AVAIL    TIMELIMIT    NODES    STATE NODELIST
nvidia*        up        5:00        7    idle apollo-nv-vm[1-7]
amd            up        5:00        1    inval apollo-amd3
amd            up        5:00        1    idle apollo-amd2
```

- ❖ Limitations
 - 1GPU or 2 GPUs per Job
 - 2 CPU cores per GPU (i.e. 1 GPU -> 2 cores, 2 GPUs -> 4 cores)
 - 2 Jobs per User
 - Wall time: 5 minutes

Instructions to compile a CUDA program

- Load cuda first!
 - `module load cuda`
- Compile
 - `nvcc -arch=sm_61 [other options] <inputfile>`
 - e.g.,
`nvcc cuda_code.cu -o cuda_executable`
 - `sm_61` is Compute Capability 6.1, which is what GTX 1080 supports
 - If you are using your own GPU, find your GPU's compute capability [here](#)
- If you have a Makefile, simply
 - `make`

Instructions to run a CUDA program

- ❖ apollo-nv-test
 - **ssh apollo-nv-test** or **ssh nv-test**
 - If you want to specify which GPU to use.
 - `export CUDA_VISIBLE_DEVICES=<gpu id>`
 - eg. `export CUDA_VISIBLE_DEVICES=1`
 - eg. `export CUDA_VISIBLE_DEVICES=0,1`
- ❖ apollo-nv-vm[1-7]
 - Slurm
 - Access gpus with the flag: `--gres=gpu:<number of gpu>`
 - eg. `srun -n 1 --gres=gpu:1 ./executable`
 - eg. `srun -n 1 --gres=gpu:2 ./executable`
 - If two GPUs are requested, they will be on the same node.

Practice

- ❖ In this practice, try to run the **deviceQuery**
- ❖ Steps:
 - cp -r /home/pp25/share/deviceQuery \$HOME
 - cd \$HOME/deviceQuery
 - nvcc deviceQuery.cpp -o deviceQuery
- ❖ Run it
 - on apollo-nv-test
 - with Slurm scheduler on apollo-nv-vm[1-7]
- ❖ How many CUDA cores on NVIDIA GTX 1080?

Device Query - Result

```
CUDA Device Query (Runtime API) version (CUDART static linking)

Detected 1 CUDA Capable device(s)

Device 0: "NVIDIA GeForce GTX 1080"
  CUDA Driver Version / Runtime Version      12.6 / 12.6
  CUDA Capability Major/Minor version number: 6.1
  Total amount of global memory:             8107 MBytes (8500871168 bytes)
  (20) Multiprocessors, (128) CUDA Cores/MP: 2560 CUDA Cores
  GPU Max Clock rate:                      1835 MHz (1.84 GHz)
  Memory Clock rate:                       5005 Mhz
  Memory Bus Width:                        256-bit
  L2 Cache Size:                          2097152 bytes
  Maximum Texture Dimension Size (x,y,z): 1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
  Maximum Layered 1D Texture Size, (num) layers: 1D=(32768), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers: 2D=(32768, 32768), 2048 layers
  Total amount of constant memory:          65536 bytes
  Total amount of shared memory per block:   49152 bytes
  Total number of registers available per block: 65536
  Warp size:                             32
  Maximum number of threads per multiprocessor: 2048
  Maximum number of threads per block:       1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size    (x,y,z): (2147483647, 65535, 65535)
  Maximum memory pitch:                  2147483647 bytes
  Texture alignment:                     512 bytes
  Concurrent copy and kernel execution: Yes with 2 copy engine(s)
  Run time limit on kernels:            No
  Integrated GPU sharing Host Memory:  No
  Support host page-locked memory mapping: Yes
  Alignment requirement for Surfaces:   Yes
  Device has ECC support:              Disabled
  Device supports Unified Addressing (UVA): Yes
  Supports Cooperative Kernel Launch: Yes
  Supports MultiDevice Co-op Kernel Launch: Yes
  Device PCI Domain ID / Bus ID / location ID: 0 / 0 / 1
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 12.6, CUDA Runtime Version = 12.6, NumDevs = 1
Result = PASS
```

Tools

nvidia-smi

- ❖ NVIDIA System Management Interface program
- ❖ You can query details about
 - gpu type
 - gpu utilization
 - memory usage
 - temperature
 - clock rate
 - ...

nvidia-smi example

```
# michael1017 @ hades02 in ~ [15:08:34]
$ nvidia-smi
Thu Nov 12 15:08:36 2020
+-----+
| NVIDIA-SMI 450.57      Driver Version: 450.57      CUDA Version: 11.0 |
+-----+
| GPU  Name      Persistence-M| Bus-Id      Disp.A | Volatile Uncorr. ECC | |
| Fan  Temp  Perf  Pwr:Usage/Cap| Memory-Usage | GPU-Util  Compute M. |
|                               |             |            | MIG M. |
+-----+
| 0  GeForce GTX 1080     On  | 00000000:4B:00.0 Off |           N/A | |
| 0%   37C    P8    7W / 200W |        1MiB /  8119MiB |     0%     Default |
|                           |             |            | N/A |
+-----+
| 1  GeForce GTX 1080     On  | 00000000:4D:00.0 Off |           N/A | |
| 0%   44C    P8    14W / 200W |        1MiB /  8117MiB |     0%     Default |
|                           |             |            | N/A |
+-----+
+-----+
| Processes:
| GPU  GI  CI      PID  Type  Process name          GPU Memory |
|           ID  ID                                         Usage
+-----+
| No running processes found
+-----+
```

compute-sanitizer

- ❖ Compute Sanitizer is a functional correctness checking suite included in the CUDA toolkit. This suite contains multiple tools that can perform different type of checks.
- ❖ Tutorial
 - [Compute-sanitizer](#)
 - [Error-type](#)
- ❖ Module
 - module load nvhpc-nompi/24.9

compute-sanitizer

```
cudaFree(device_t);
cudaFree(device_t); // free an address twice, error
```

```
gilbert12@apollo-nv-test:~/lab-sobel$ compute-sanitizer ./sobel testcases/candy.png candy.out.png
===== COMPUTE-SANITIZER
===== Program hit cudaErrorInvalidValue (error 1) due to "invalid argument" on CUDA API call to cudaFree.
=====     Saved host backtrace up to driver entry point at error
=====     Host Frame: [0x419b65]
=====             in /lib/x86_64-linux-gnu/libcuda.so.1
=====     Host Frame:cudaFree [0x573c7]
=====             in /home/gilbert12/lab-sobel./sobel
=====     Host Frame:main [0xacde]
=====             in /home/gilbert12/lab-sobel./sobel
=====     Host Frame:_libc_start_call_main in ../sysdeps/nptl/libc_start_call_main.h:58 [0x27249]
=====             in /lib/x86_64-linux-gnu/libc.so.6
=====     Host Frame:_libc_start_main in ../../csu/libc-start.c:360 [0x27304]
=====             in /lib/x86_64-linux-gnu/libc.so.6
=====     Host Frame:_start [0xaf80]
=====             in /home/gilbert12/lab-sobel./sobel
=====
===== ERROR SUMMARY: 1 error
```

cuda-gdb

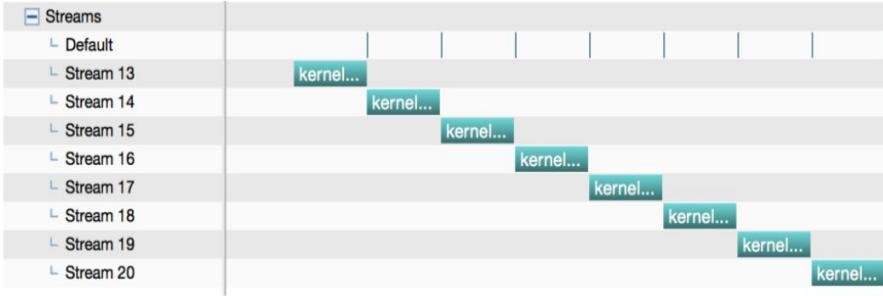
- ❖ [cuda-gdb tutorial](#)
- ❖ Module
 - module load nvhpc-nompi/24.9

nvprof

- ❖ nvprof provide you feedback about how to optimize CUDA programs
 - nvprof <CUDA executable>
 - -o <FILE> to save result to a file
 - -i <FILE> to read result from a file

nvvp

- ❖ [nvvp-tutorial](#)
- ❖ GUI version of nvprof
- ❖ Useful for the stream optimization
 - Timeline



nvvp is useful for checking the concurrency of stream

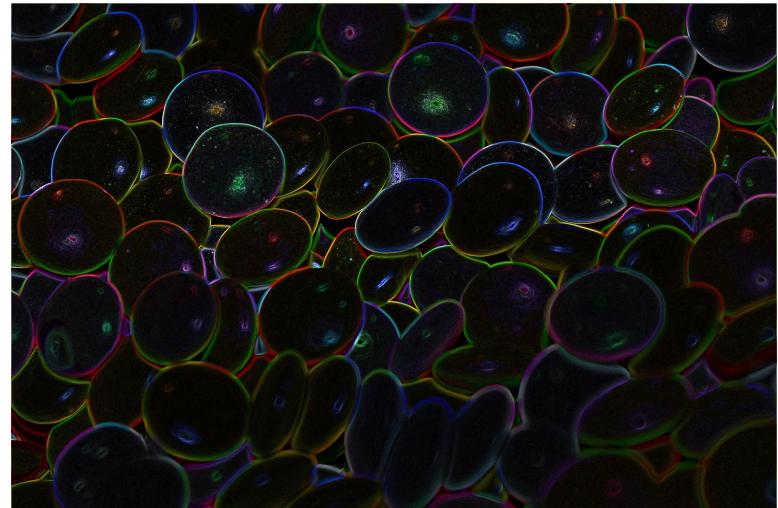
NSight Systems (nsys)

- ❖ [nsys-tutorial](#)
- ❖ Module
 - module load nvhpc-nompi/24.9
- ❖ Usage
 - nsys profile -t cuda ...

Lab3 Assignment

Problem Description

- ❖ Edge Detection: Identifying points in a digital image at which the image brightness changes sharply



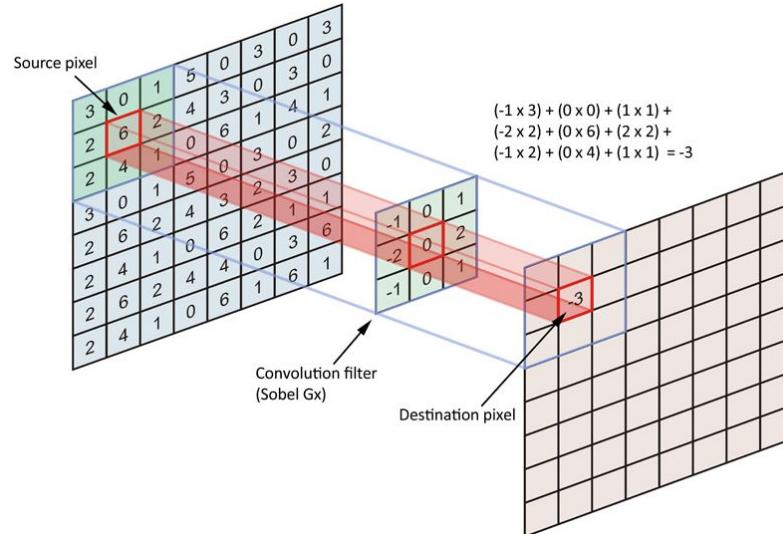
Sobel Operator

- ❖ Used in image processing and computer vision, particularly within **edge detection algorithms**.
- ❖ Uses two **3x3 filter matrix g_x, g_y** which are **convolved with the original image** to calculate approximations of the derivatives - one for horizontal changes, and one for vertical.
- ❖ In this lab, we use **5x5 kernels**

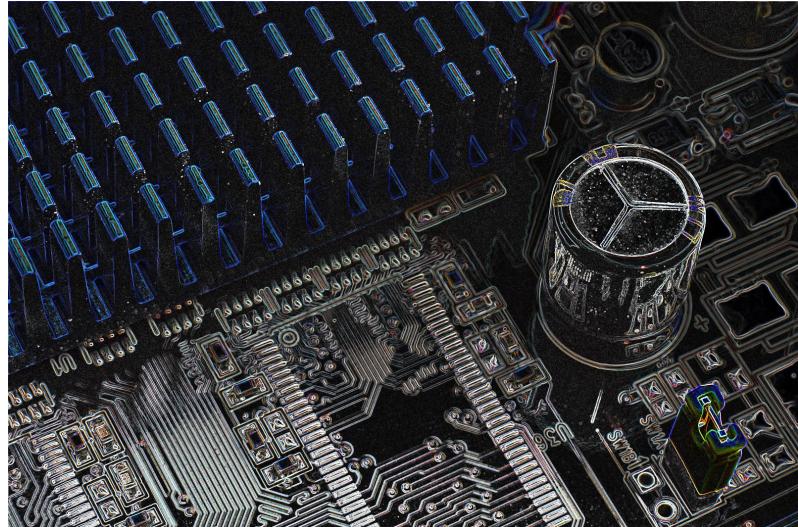
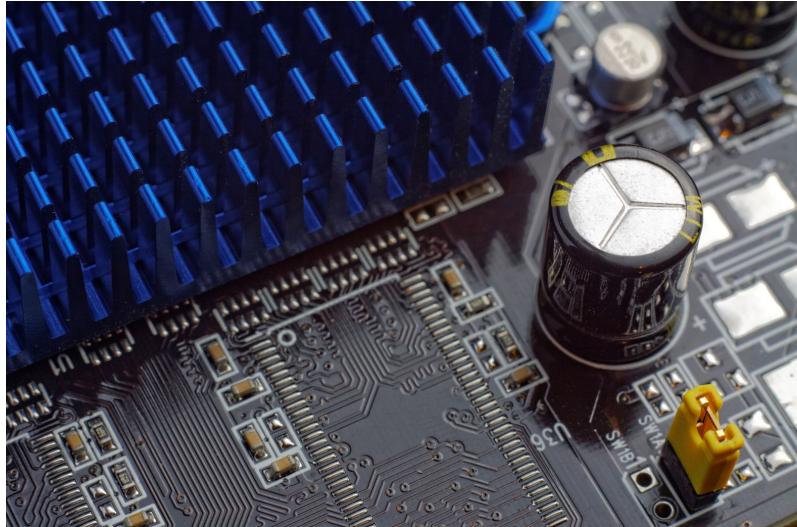
$$g_x = \begin{pmatrix} -1 & -2 & 0 & 2 & 1 \\ -4 & -8 & 0 & 8 & 4 \\ -6 & -12 & 0 & 6 & 12 \\ -4 & -8 & 0 & 8 & 4 \\ -1 & -2 & 0 & 2 & 1 \end{pmatrix}, \quad g_y = \begin{pmatrix} -1 & -4 & -6 & -4 & -1 \\ -2 & -8 & -12 & -8 & -2 \\ 0 & 0 & 0 & 0 & 0 \\ 2 & 8 & 12 & 8 & 2 \\ 1 & 4 & 6 & 4 & 1 \end{pmatrix}$$

Convolution Calculation

- ❖ Iterate through the width and height of the image
- ❖ For each pixel, multiply the filter matrix with original image element-wisely and sum them up.



Sample Result



Preparation

- ❖ TA provided CPU version, Makefile, and hint
- ❖ Files are located at /home/pp25/share/lab-sobel
- ❖ Please do not modify the test cases
- ❖ sobel.cu is cpu version (you need to rewrite it with cuda!)
- ❖ Follow hints
- ❖ After you finish the code using CUDA, try it on AMD GPU

Workflow

1. `cp -r /home/pp25/share/lab-sobel $HOME`
2. `module load rocm cuda`
3. Finish the hints
4. Compile the program : `make sobel`
5. Run the program : `srun --gres=gpu:1 ./sobel testcases/candy.png candy.out.png`
6. Check the diff : `png-diff testcases/candy.out.png candy.out.png`
7. Judge: `lab-sobel-judge`
8. Scoreboard: [sobel](#)

Workflow for AMD GPU (Optional)

We ask you to test your code on AMD GPU as well.

No need to rewrite the code, just use “Hipify”

1. `module load rocm cuda`
2. Hipify your CUDA code: `hipify-clang sobel.cu`
Generates `sobel.cu.hip`
3. Inspect the code and learn how HIP works
4. Rename the file to `sobel.hip`
5. Compile the program
`make sobel-amd`
6. Run : `srun -p amd ./sobel-amd testcases-amd/candy.png candy.out.png`
7. Judge: `lab-sobel-amd-judge`
8. Scoreboard: `sobel-amd`

How to run

- ❖ **apollo-nv-test**
 - `./sobel <input> <output>`
 - `CUDA_VISIBLE_DEVICES=0 ./sobel <input> <output>`
- ❖ **apollo-nv-vm[1-7]**
 - `srun -n 1 --gres=gpu:1 ./sobel <input> <output>`
- ❖ **apollo-amd-vm**
 - `srun -p amd -n 1 --gres=gpu:1 ./sobel-hip <input> <output>`

Check the correctness

- ❖ `png-diff <result_file> <answer_file>`
 - It verifies the correctness of your output result
 - `result_file` is the output file from your CUDA program.
 - `answer_file` is the provided file for correctness checking.
 - If your `input_file` is “`~/lab_sobel/testcases/candy.png`”,
your `answer_file` is “`~/lab_sobel/testcases/candy.out.png`”

```
[kswang@hades02 lab]$ png-diff testcases/candy.out.png test.png
ok, 100.00% 😊
```
 - Your code is correct if you see “`ok, 100.00%`”

Hints

- ❖ Malloc memory on GPU
- ❖ Copy the original image to GPU
- ❖ Put filter matrix on device memory (or declare it on device)
- ❖ Parallelize the sobel computing
- ❖ Copy the results from device to host
- ❖ Free unused address

Submission

- Judge will execute your code with single process, single GPU
- Submit your code and Makefile (optional) to eeclass before
11/6 23:59
- Use **lab-sobel-judge** to judge
- Score Board: [lab-sobel](#)
- Get started as soon as possible to avoid heavy queueing delay
 - sobel.cu
 - Makefile (Optional)