

# OpenACC

Parallel Programming  
2025/10/30



國立清華大學叢集電腦競賽團隊  
Student Cluster Competition Team of NTHU

# What is OpenACC?

- **Open Accelerators**
- Through various **compiler directives** to write GPU code
- Lower the technical barriers to GPU programming



# What is OpenACC

```
#pragma acc data copy(A) create(Anew)
while ( error > tol && iter < iter_max ) {
    error = 0.0;
#pragma acc kernels
{
#pragma acc loop independent collapse(2) reduction(max:error)
    for ( int j = 1; j < n-1; j++ ) {
        for ( int i = 1; i < m-1; i++ ) {
            Anew [j] [i] = 0.25 * ( A [j] [i+1] + A [j] [i-1] +
                                      A [j-1] [i] + A [j+1] [i]);
            error = max ( error, fabs (Anew [j] [i] - A [j] [i]));
        }
    }
    ...
}
```

<https://www.openacc.org/>



# CUDA

- CudaMalloc(...): Declare memory on the GPU
- CudaMemcpy(...): Move data
- functionname<<<thread, blocks>>>(...): Write your own Cuda

Kernel Function

=> High entry barrier



# CUDA Example

```
#define N 10000000
int main() {
    float x[N], y[N];
    for (int i = 0; i < N; i++) {
        y[i] = sin(x[i]);
    }
}
```



```
#define N 10000000
// CUDA kernel function
__global__ void cuda_sin(float *x, float *y, int n) {
    int threadId = blockIdx.x * blockDim.x + threadIdx.x;
    if (threadId >= n) {
        return;
    }
    y[threadId] = sin(x[threadId]);
}

int main() {
    float x[N], y[N];
    // 告訴 GPU device 上的 array
    float *x_device, *y_device;
    cudaMalloc(&x_device, sizeof(float) * N);
    cudaMalloc(&y_device, sizeof(float) * N);
    // 將 x 陣列複製到 GPU 上
    cudaMemcpy(x_device, x, sizeof(float) * N, cudaMemcpyHostToDevice);
    // 計算所需要的 thread block 數量
    int thread_block_sz = 1024;
    int num_blocks = N / thread_block_sz;
    if (N % thread_block_sz != 0) {
        num_blocks += 1;
    }
    // 呼叫 CUDA kernel function
    cuda_sin<<<num_blocks, thread_block_sz>>>(x_device, y_device, N);
    // 把結果複製回 y 陣列
    cudaMemcpy(y, y_device, sizeof(float) * N, cudaMemcpyDeviceToHost);
    // 釋放 GPU 記憶體
    cudaFree(x_device);
    cudaFree(y_device);
}
```



# OpenACC

- No need to declare the memory on the device
- `#pragma acc data copy(...)`: You can move data with a simple clause
- You can directly use parallel region to port to the GPU.

=> Easy to use

```
#define N 10000000
int main() {
    float x[N], y[N];
    #pragma acc data copyin(x[0:N]) copyout(y[0:N])
    #pragma acc parallel loop
    for (int i = 0; i < N; i++) {
        y[i] = sin(x[i]);
    }
}
```



# OpenACC Directive

- `#pragma acc <directive> <clause>`
  - `#pragma` is a compiler hint
  - `acc` tells the compiler that this is the OpenACC pragma
  - `directive` is what OpenACC tells the compiler to indicate
  - `clause` is an instruction for OpenACC to supplement or optimize the directive.
  - will take effect on next `block/line`.

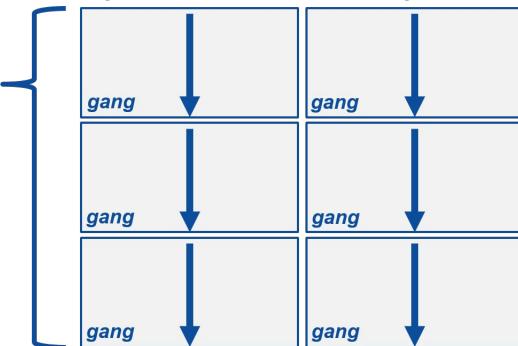


# OpenACC Directive - parallel

- `#pragma acc parallel`
  - parallel tells the compiler that this code should be **redundantly parallelized**
  - tells the compiler to run this part of code on **GPU**

```
#pragma acc parallel {
```

When encountering the **parallel** directive, the compiler will generate 1 or more parallel **gangs**, which execute redundantly.



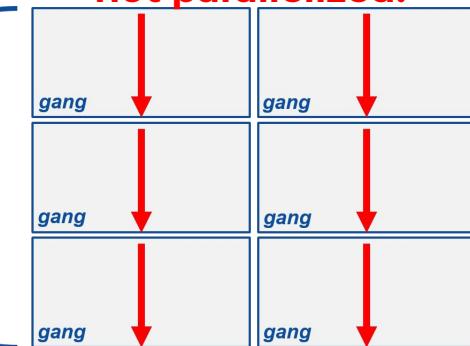
```
#pragma acc parallel {
```

```
for(int i = 0; i < N; i++)  
{  
    // Do Something  
}
```

This loop will be **redundantly parallelized** across the **gangs**

This means that each **gang** will execute the entire loop

**not parallelized!**



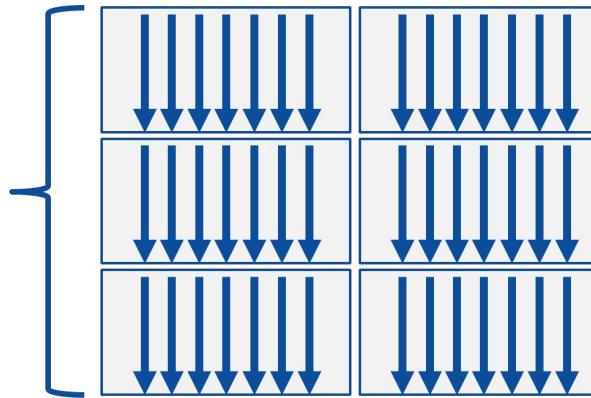
# OpenACC Directive - loop

- `#pragma acc parallel loop`
  - `loop` tells the compiler that this loop needs to be parallelized
  - It also tells the compiler that this loop can be safely parallelized.

```
#pragma acc parallel  
{  
  
    #pragma acc loop  
    for(int i = 0; i < N; i++)  
    {  
        // Do Something  
    }  
}
```

The `loop` directive informs the compiler which loops to parallelize.

The iterations of the loop will be broken up evenly among the parallel `gangs`.



The `gangs` will then execute in parallel with one another.



# OpenACC Directive - reduction

- `#pragma acc parallel loop reduction(<operation>:<target>)`
  - `reduction` tells the compiler that a target is to be reduced
  - `reduce`: perform global operations on the selected target

```
int sum = 0;  
#pragma acc parallel loop reduction(+:sum)  
for(int i = 0, i < N, i++) sum += i;
```

reduction(+ : sum)

reduction(\* : prod)

reduction(min : mn)

reduction(& : mask)



# OpenACC Directive - reduction cont'd

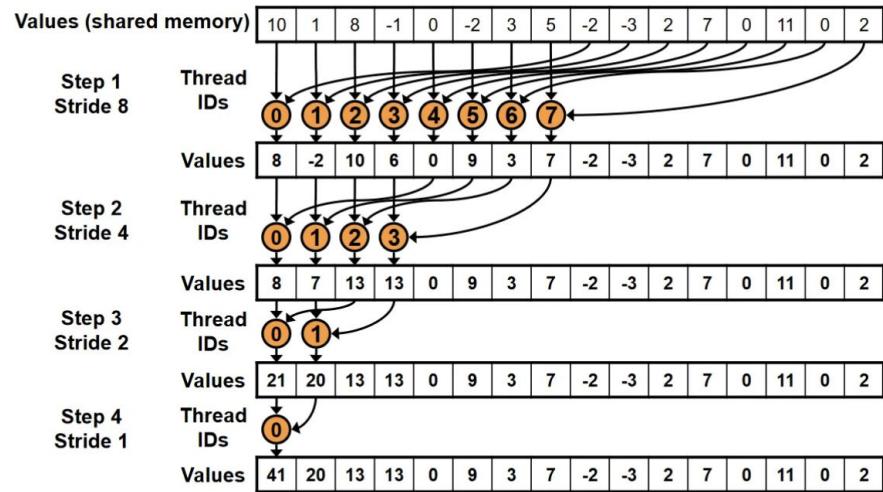
- How reduction work?
  - Each gang computes its own local result in parallel.
  - After the loop ends, OpenACC reduces all gang results into a single final value.

```
int main() {
    float sum = 0;
    float mn = 999999;
    float x[N];
    #pragma acc parallel loop reduction(min:mn) reduction(+:sum)
    for (int i = 0; i < N; i++) {
        float tmp = sin(x[i]);
        sum = sum + tmp;
        mn = min(mn, tmp);
    }
}
```



# OpenACC Directive - parallel reduction

- Works only when there's no data dependency between iterations.
- Reduces in a tree-like fashion — pairs of partial results are combined.
- Overall time complexity becomes  $O(\log N)$ .  
( $N$  is the number of gangs)



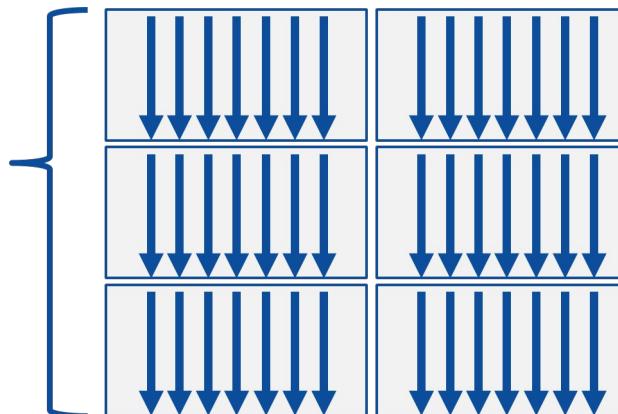
# OpenACC Directive - kernels

- **#pragma acc kernels**
  - All actions are decided by the compiler
  - You can also include the sequential code

```
#pragma acc kernels
{
    for(int i = 0; i < N; i++)
    {
        // Do Something
    }

    for(int i = 0; i < M; i++)
    {
        // Do Something Else
    }
}
```

With the ***kernel*** directive, the ***loop*** directive is implied.



# OpenACC Directive - kernels

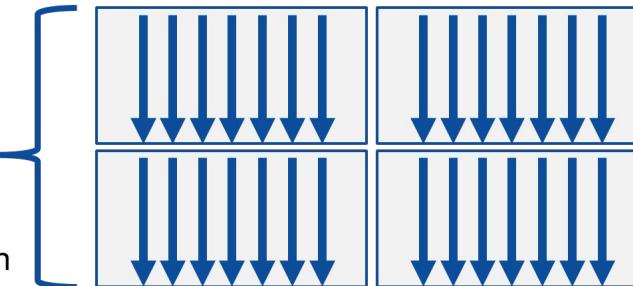
- `#pragma acc kernels loop independent`
  - Tell the compiler that this loop can be safely parallelized, and force parallelization of it

```
#pragma acc kernels
{
    for(int i = 0; i < N; i++)
    {
        // Do Something
    }

    for(int i = 0; i < M; i++)
    {
        // Do Something Else
    }
}
```

This process can happen multiple times within the **kernels** region.

Each loop can have a different number of gangs, and those gangs can be organized/optimized completely differently.



```
#pragma acc kernels
{
    for (int i = 0; i < N; i++)
    {
        < loop code >
    }

    < some other sequential code >

    #pragma acc loop independent
    for (int j = 0; j < M; j++)
    {
        < loop code >
    }
}
```



# Data Management

```
#pragma acc data copy(A[0:N])
```

```
#pragma acc parallel
```

```
{
```

```
#pragma acc loop
```

```
for(int i = 0; i < N; i++) A[i] = 0;
```

```
}
```



```
#pragma acc kernels copy(a[0:N])
for(int i = 0; i < N; i++){
    a[i] = 0;
}
```



# Data Management

You can copy only part of the data  
#pragma acc parallel loop copy(A[1:N-2])

- **#pragma acc data copy(...)**
  - Copy the data into the GPU and copy the data back to the CPU after the parallel region ends
- **#pragma acc data copyin(...)**
  - Copy the data into the GPU and delete the data on the GPU after the parallel region ends.
- **#pragma acc data copyout(...)**
  - Copy the data back to the CPU and delete the data on the GPU after the parallel region ends.
- **#pragma acc data create(...)**
  - Declare a space on the GPU without performing any copying operations
  - When there are variables for temporary storage, using this clause eliminates the need to copy in and out.



# Loop Optimization

- `#pragma acc parallel loop collapse(...)`
  - Can be used in tightly nested loops
  - collapse can flatten loops and turn multiple loops into one large parallel loop

```
#pragma acc parallel loop collapse( 2 )
```

```
for(int j = 0; j < M; j++) {  
    for(int k = 0; k < Q; k++) {  
        < loop code >  
    }  
}
```

## TIP1:

When the outer loop is too small, flattening the loop can increase GPU usage.



# Loop Optimization

- `#pragma acc parallel loop tile(x, y)`
  - Calculate **nested** loop break for multiple tiles (blocks)

```
#pragma acc parallel loop tile( 32, 32 )
```

```
for(int j = 0; j < 128; j++) {  
    for(int k = 0; k < 128; k++) {  
        < loop code >  
    }  
}
```

## TIP1:

Try to make the tile size a multiple of 32. The threads in a worker and vector of Nvidia GPU are executed in units of 32.

## TIP2:

Do not use tiles larger than 32\*32, because in NVIDIA GPU, the maximum number of threads in a gang is 1024 (32\*32)



# Practice

Try to parallelize this code.

```
#include <bits/stdc++.h>
#define N 100000
double x[N], y[N], sum = 0;
void generate_random() {
    srand(0);
    for (int i = 0; i < N; i++) {
        x[i] = (rand() / (double)RAND_MAX) * 10.0 - 5.0;
    }
}

int main() {
    generate_random();
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            y[i] += x[i] * x[j] + cos(x[i] + x[j]);
            sum = sum + y[i];
        }
    }
    std::cout << sum << '\n';
}
```

Execution time >30 seconds



# How to compile

```
nvc++ -fast -acc -gpu=cc61 -Minfo=accel -o <執行檔名稱> <檔案.cpp>
```

- **-acc** tell compiler to support OpenACC
- **-gpu=cc61** assign GPU architecture GTX1080(cc61), V100(cc70)
- **-Minfo=accel** print info about the code is parallelized

```
gilbert12@apollo-login:~$ nvc++ -fast -acc -gpu=cc61 -Minfo=accel -o acc acc.cpp
main:
  12, Generating copyin(x[:]) [if not already present]
    Generating NVIDIA GPU code
    15, #pragma acc loop gang, vector(128) /* blockIdx.x threadIdx.x */
      Generating reduction(+:sum)
    16, #pragma acc loop seq
    12, Generating copyout(y[:]) [if not already present]
      Generating implicit copy(sum) [if not already present]
    16, Complex loop carried dependence of y prevents parallelization
      Loop carried dependence of y prevents parallelization
      Loop carried backward dependence of y prevents vectorization
```



# Answer

```
#pragma acc data copyin(x[0:N]) copyout(y[0:N])
#pragma acc parallel loop reduction(+:sum)
for (int i = 0; i < N; i++) {
    for (int j = 0; j < N; j++) {
        y[i] += x[i] * x[j] + cos(x[i] + x[j]);
        sum = sum + y[i];
    }
}
std::cout << sum << '\n';
```

Execution time < 1 second



# OpenACC Labs



# Deep Neural Network(DNN)

- MNIST hand written digits classification Inference

label = 5



label = 0



label = 4



label = 1



label = 9



label = 2



label = 1



label = 3



label = 1



label = 4



label = 3



label = 5



label = 3



label = 6



label = 1



label = 7



label = 2



label = 8



label = 6



label = 9



# Deep Neural Network(DNN)

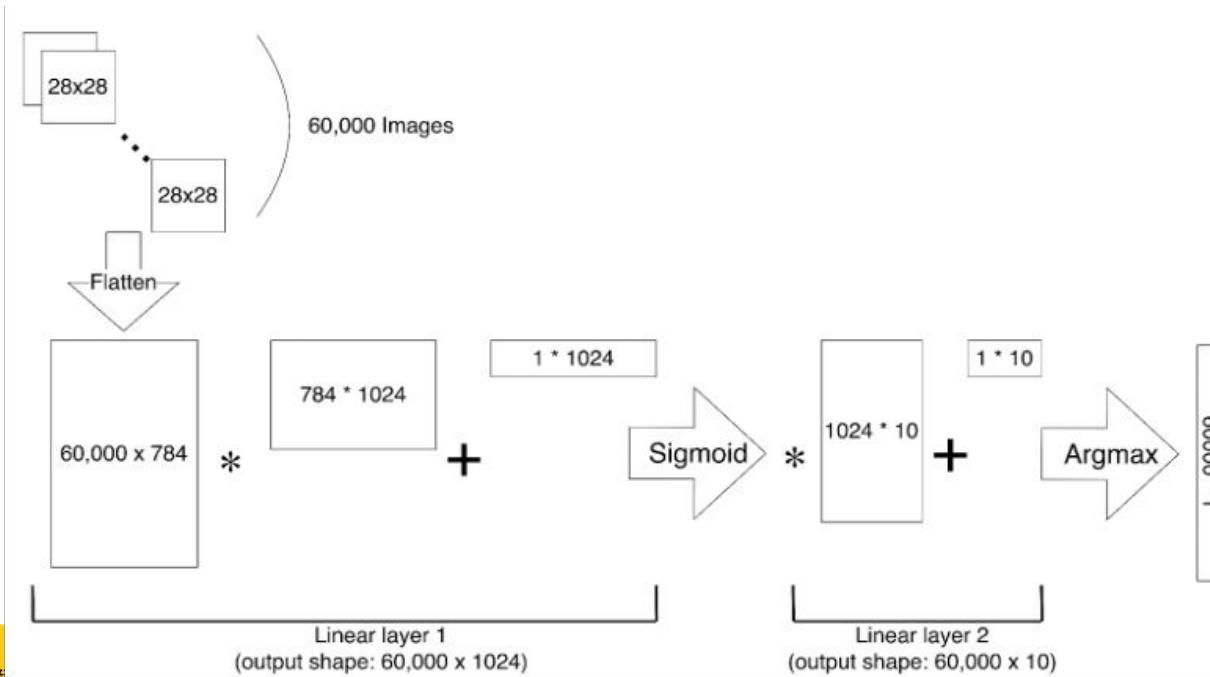
- Two fully-connected layer

```
simpleNN(  
    nn): Sequential(  
        (0): Flatten(start_dim=1, end_dim=-1)  
        (1): Linear(in_features=784, out_features=1024, bias=True)  
        (2): Sigmoid()  
        (3): Linear(in_features=1024, out_features=10, bias=True)  
        (4): Softmax(dim=None)  
    )  
)
```



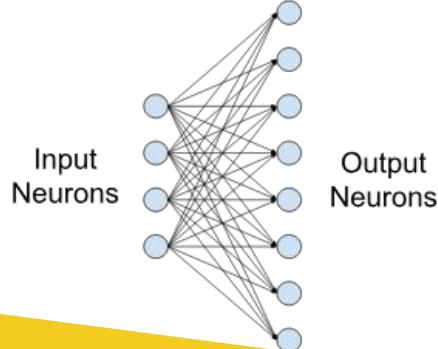
# Deep Neural Network(DNN)

- Two fully-connected layer

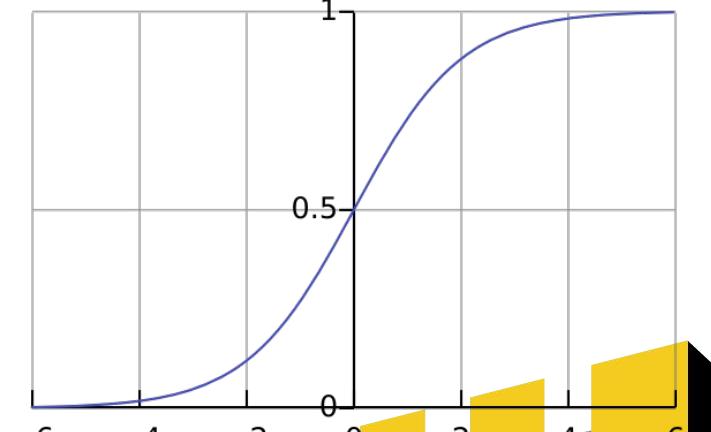


# Deep Neural Network(DNN)

- Two fully-connected layer
  - First Linear Layer + Sigmoid Activation Function
  - Second Linear Layer + Argmax



$$\text{Sigmoid}(x) = \sigma(x) = \frac{1}{1 + \exp(-x)}$$



# Task

- Given the sequential version of DNN code that runs on CPU.
- Try to parallel the matrix computation in the neural network.
- Using OpenACC



# Task

- Provided files:
  - Sequential code
  - pretrained model weights(since the task is inference)
  - Makefile
- Please only modify the functions with TODO label.



# Workflow

1. `cp -r /home/pp25/share/lab-mnist $HOME`
2. `cd ~/lab-mnist`
3. Parallel the TODOs in `mnist.cpp`
4. `module load nvhpc-nompi`
5. compile the program : `make`
6. run the DNN program : `srun -p nvidia --gres=gpu:1 ./mnist`
7. judge: `lab-mnist-judge`  
scoreboard: [lab-mnist](#)



# Result

- Inference accuracy should be **97.8183%** (the parallelization should not affect the accuracy)



# Result

```
gilbert12@apollo-login:~/lab-mnist$ srun -p nvidia mnist
srun: No GPU specified, defaulting to allocating 1 GPU
CUDA initialized.

Nbr of training images = 60000
Reading file: /share/testcases/lab-mnist/weights/layer1_matrix
Reading file: /share/testcases/lab-mnist/weights/layer1_bias
Reading file: /share/testcases/lab-mnist/weights/layer2_matrix
Reading file: /share/testcases/lab-mnist/weights/layer2_bias

Inference accuracy: 97.8183%

----- STATS -----
Time for initializing CUDA device:      104 m.s.
Time for reading MNIST data & weights: 173 m.s.
Time for inferencing                  : 545 m.s.
Time for calculating accuracy        : 6 m.s.
----- END OF STATS -----
```



# Judge Result

- Judge: lab-mnist-judge
- Scoreboard: [lab-mnist](#)
- Platform: apollo-gpu

