

# Lab6 FlashAttention

Nov, 2025 Parallel Programming

# Outline

- Attention
- FlashAttention
- Lab6 Assignment
- HW4 Assignment

# Attention

- Q : What we're focusing on.
- K : What features are available.
- V : What content is retrieved based on focus.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \quad \mathbf{Q}, \mathbf{K}, \mathbf{V} \in \mathbb{R}^{N \times d}$$

# Attention

$$\mathbf{S} = \mathbf{Q}\mathbf{K}^\top \in \mathbb{R}^{N \times N}, \quad \mathbf{P} = \text{softmax}(\mathbf{S}) \in \mathbb{R}^{N \times N}, \quad \mathbf{O} = \mathbf{P}\mathbf{V} \in \mathbb{R}^{N \times d},$$

$$\begin{bmatrix} s_{11} & s_{12} & s_{13} & s_{14} \\ s_{21} & s_{22} & s_{23} & s_{24} \\ s_{31} & s_{32} & s_{33} & s_{34} \\ s_{41} & s_{42} & s_{43} & s_{44} \end{bmatrix} = \begin{bmatrix} q_{11} & q_{12} & q_{13} \\ q_{21} & q_{22} & q_{23} \\ q_{31} & q_{32} & q_{33} \\ q_{41} & q_{42} & q_{43} \end{bmatrix} \cdot \begin{bmatrix} k_{11} & k_{21} & k_{31} & k_{41} \\ k_{12} & k_{22} & k_{32} & k_{42} \\ k_{13} & k_{23} & k_{33} & k_{43} \end{bmatrix}$$

# Attention

$$\mathbf{S} = \mathbf{Q}\mathbf{K}^\top \in \mathbb{R}^{N \times N}, \quad \boxed{\mathbf{P} = \text{softmax}(\mathbf{S}) \in \mathbb{R}^{N \times N}}, \quad \mathbf{O} = \mathbf{P}\mathbf{V} \in \mathbb{R}^{N \times d},$$

$$\begin{bmatrix} p_{11} & p_{12} & p_{13} & p_{14} \\ p_{21} & p_{22} & p_{23} & p_{24} \\ p_{31} & p_{32} & p_{33} & p_{34} \\ p_{41} & p_{42} & p_{43} & p_{44} \end{bmatrix} = \text{softmax} \left( \begin{bmatrix} s_{11} & s_{12} & s_{13} & s_{14} \\ s_{21} & s_{22} & s_{23} & s_{24} \\ s_{31} & s_{32} & s_{33} & s_{34} \\ s_{41} & s_{42} & s_{43} & s_{44} \end{bmatrix} \right)$$

# Attention

$$\mathbf{S} = \mathbf{Q}\mathbf{K}^\top \in \mathbb{R}^{N \times N}, \quad \mathbf{P} = \text{softmax}(\mathbf{S}) \in \mathbb{R}^{N \times N}, \quad \boxed{\mathbf{O} = \mathbf{P}\mathbf{V} \in \mathbb{R}^{N \times d}},$$

$$\begin{bmatrix} o_{11} & o_{12} & o_{13} \\ o_{21} & o_{22} & o_{23} \\ o_{31} & o_{32} & o_{33} \\ o_{41} & o_{42} & o_{43} \end{bmatrix} = \begin{bmatrix} p_{11} & p_{12} & p_{13} & p_{14} \\ p_{21} & p_{22} & p_{23} & p_{24} \\ p_{31} & p_{32} & p_{33} & p_{34} \\ p_{41} & p_{42} & p_{43} & p_{44} \end{bmatrix} \cdot \begin{bmatrix} v_{11} & v_{12} & v_{13} \\ v_{21} & v_{22} & v_{23} \\ v_{31} & v_{32} & v_{33} \\ v_{41} & v_{42} & v_{43} \end{bmatrix}.$$

# Multi-Head Attention

- Rich Representations
- Efficient Parallelization
- E.g. `emb_dim = 4096`  $\rightarrow$  `num_heads = 32, head_size = 128`

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$

where  $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$

# Causal Attention

- If you're predicting the next word in a sentence, the model shouldn't have access to future words beyond the current position.

$s_{11}$	$s_{12}$	$s_{13}$	$s_{14}$
$s_{21}$	$s_{22}$	$s_{23}$	$s_{24}$
$s_{31}$	$s_{32}$	$s_{33}$	$s_{34}$
$s_{41}$	$s_{42}$	$s_{43}$	$s_{44}$

$$\text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

$$\text{Masked Scores}_{i,j} = \begin{cases} \frac{(QK^T)_{i,j}}{\sqrt{d_k}}, & \text{if } j \leq i \\ -\infty, & \text{if } j > i \end{cases}$$

$$\text{Masked Attention Weights}_{i,j} = \text{softmax}(\text{Masked Scores}_{i,j})$$

$$\text{Causal Attention}(Q, K, V) = \text{softmax} \left( \text{Mask} \left( \frac{QK^T}{\sqrt{d_k}} \right) \right) V$$

# FlashAttention - Overview

- Goal: avoid reading and writing the attention matrix to and from HBM.
  - Computing the softmax without access to the whole input.
  - Not storing the large intermediate attention matrix for the backward pass.

---

## Algorithm 0 Standard Attention Implementation

---

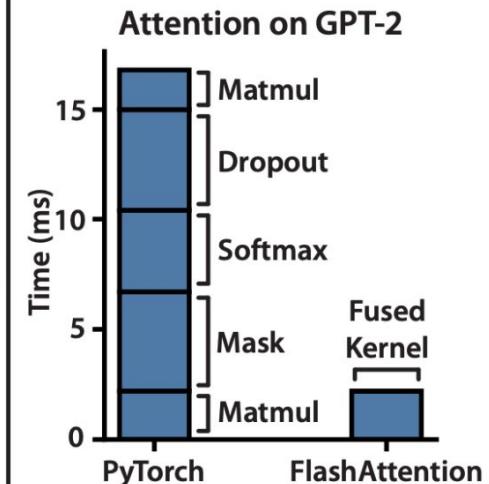
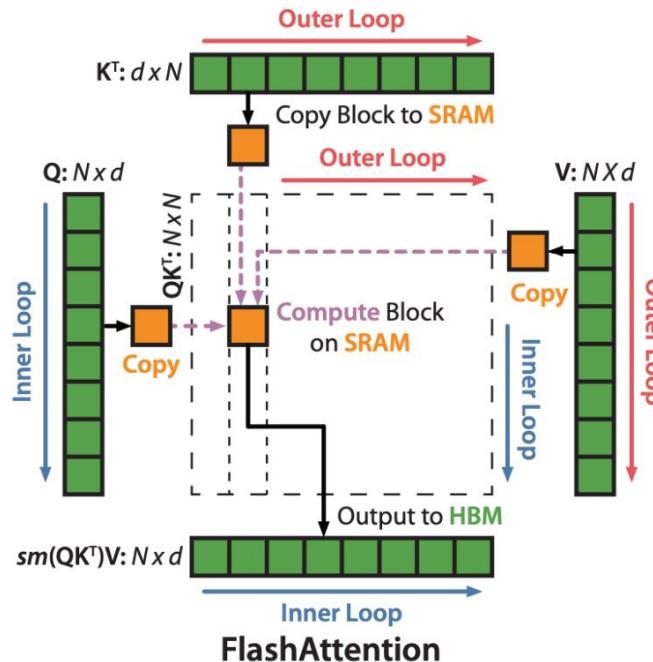
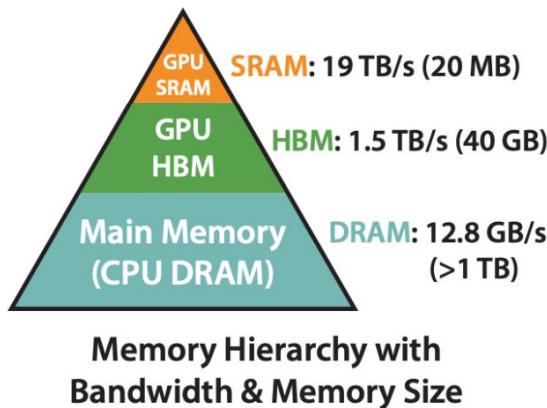
**Require:** Matrices  $\mathbf{Q}, \mathbf{K}, \mathbf{V} \in \mathbb{R}^{N \times d}$  in HBM.

- 1: Load  $\mathbf{Q}, \mathbf{K}$  by blocks from HBM, compute  $\mathbf{S} = \mathbf{Q}\mathbf{K}^\top$ , write  $\mathbf{S}$  to HBM.
  - 2: Read  $\mathbf{S}$  from HBM, compute  $\mathbf{P} = \text{softmax}(\mathbf{S})$ , write  $\mathbf{P}$  to HBM.
  - 3: Load  $\mathbf{P}$  and  $\mathbf{V}$  by blocks from HBM, compute  $\mathbf{O} = \mathbf{P}\mathbf{V}$ , write  $\mathbf{O}$  to HBM.
  - 4: Return  $\mathbf{O}$ .
-

# FlashAttention - Method

- Tiling: split the input into blocks and make several passes over input blocks.
  - Matrix multiplication and pointwise operations are easy to handle.
  - SoftMax: need to maintain  $m(x)$ ,  $l(x)$ .
- Recompute: store the softmax normalization factor in order to quickly recompute in the backward pass.

# FlashAttention



# FlashAttention - SoftMax

$$\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_{j=1}^d e^{x_j}}$$

$$m = \max_i(x_i); \quad \text{softmax}(x_i) = \frac{e^{x_i-m}}{\sum_{j=1}^d e^{x_j-m}}$$

# FlashAttention - SoftMax

$$m(x) := \max_i x_i, \quad f(x) := [e^{x_1-m(x)} \quad \dots \quad e^{x_B-m(x)}], \quad \ell(x) := \sum_i f(x)_i, \quad \text{softmax}(x) := \frac{f(x)}{\ell(x)}.$$

For vectors  $x^{(1)}, x^{(2)} \in \mathbb{R}^B$ , we can decompose the softmax of the concatenated  $x = [x^{(1)} \ x^{(2)}] \in \mathbb{R}^{2B}$  as:

$$\begin{aligned} m(x) &= m([x^{(1)} \ x^{(2)}]) = \max(m(x^{(1)}), m(x^{(2)})), \quad f(x) = [e^{m(x^{(1)})-m(x)} f(x^{(1)}) \quad e^{m(x^{(2)})-m(x)} f(x^{(2)})], \\ \ell(x) &= \ell([x^{(1)} \ x^{(2)}]) = e^{m(x^{(1)})-m(x)} \ell(x^{(1)}) + e^{m(x^{(2)})-m(x)} \ell(x^{(2)}), \quad \text{softmax}(x) = \frac{f(x)}{\ell(x)}. \end{aligned}$$

# FlashAttention - SoftMax

$$m_1 = \max([1, 2]) = 2$$

$$m_2 = \max([3, 4]) = 4$$

$$m = \max(m_1, m_2) = 4$$

$$f_1 = [e^{1-2}, e^{2-2}] = [e^{-1}, e^0] \quad f_2 = [e^{3-4}, e^{4-4}] = [e^{-1}, e^0] \quad f = [e^{m_1-m} f_1, e^{m_2-m} f_2] = [e^{-3}, e^{-2}, e^{-1}, e^0]$$

$$l_1 = \sum f_1 = e^{-1} + e^0 \quad l_2 = \sum f_2 = e^{-1} + e^0 \quad l = e^{m_1-m} l_1 + e^{m_2-m} l_2 = e^{-3} + e^{-2} + e^{-1} + e^0$$

$$o_1 = \frac{f_1}{l_1} = \frac{[e^{-1}, e^0]}{e^{-1} + e^0}$$

$$o_2 = \frac{f_2}{l_2} = \frac{[e^{-1}, e^0]}{e^{-1} + e^0}$$

$$o = \frac{f}{l} = \frac{[e^{-3}, e^{-2}, e^{-1}, e^0]}{e^{-3} + e^{-2} + e^{-1} + e^0}$$

# FlashAttention - Algorithm

---

**Algorithm 1** FLASHATTENTION

---

**Require:** Matrices  $\mathbf{Q}, \mathbf{K}, \mathbf{V} \in \mathbb{R}^{N \times d}$  in HBM, on-chip SRAM of size  $M$ .

- 1: Set block sizes  $B_c = \lceil \frac{M}{4d} \rceil, B_r = \min(\lceil \frac{M}{4d} \rceil, d)$ .
  - 2: Initialize  $\mathbf{O} = (0)_{N \times d} \in \mathbb{R}^{N \times d}, \ell = (0)_N \in \mathbb{R}^N, m = (-\infty)_N \in \mathbb{R}^N$  in HBM.
  - 3: Divide  $\mathbf{Q}$  into  $T_r = \lceil \frac{N}{B_r} \rceil$  blocks  $\mathbf{Q}_1, \dots, \mathbf{Q}_{T_r}$  of size  $B_r \times d$  each, and divide  $\mathbf{K}, \mathbf{V}$  in to  $T_c = \lceil \frac{N}{B_c} \rceil$  blocks  $\mathbf{K}_1, \dots, \mathbf{K}_{T_c}$  and  $\mathbf{V}_1, \dots, \mathbf{V}_{T_c}$ , of size  $B_c \times d$  each.
  - 4: Divide  $\mathbf{O}$  into  $T_r$  blocks  $\mathbf{O}_1, \dots, \mathbf{O}_{T_r}$  of size  $B_r \times d$  each, divide  $\ell$  into  $T_r$  blocks  $\ell_1, \dots, \ell_{T_r}$  of size  $B_r$  each, divide  $m$  into  $T_r$  blocks  $m_1, \dots, m_{T_r}$  of size  $B_r$  each.
  - 5: **for**  $1 \leq j \leq T_c$  **do**
  - 6:   Load  $\mathbf{K}_j, \mathbf{V}_j$  from HBM to on-chip SRAM.
  - 7:   **for**  $1 \leq i \leq T_r$  **do**
  - 8:     Load  $\mathbf{Q}_i, \mathbf{O}_i, \ell_i, m_i$  from HBM to on-chip SRAM.
  - 9:     On chip, compute  $\mathbf{S}_{ij} = \mathbf{Q}_i \mathbf{K}_j^T \in \mathbb{R}^{B_r \times B_c}$ .
  - 10:    On chip, compute  $\tilde{m}_{ij} = \text{rowmax}(\mathbf{S}_{ij}) \in \mathbb{R}^{B_r}, \tilde{\mathbf{P}}_{ij} = \exp(\mathbf{S}_{ij} - \tilde{m}_{ij}) \in \mathbb{R}^{B_r \times B_c}$  (pointwise),  $\tilde{\ell}_{ij} = \text{rowsum}(\tilde{\mathbf{P}}_{ij}) \in \mathbb{R}^{B_r}$ .
  - 11:    On chip, compute  $m_i^{\text{new}} = \max(m_i, \tilde{m}_{ij}) \in \mathbb{R}^{B_r}, \ell_i^{\text{new}} = e^{m_i - m_i^{\text{new}}} \ell_i + e^{\tilde{m}_{ij} - m_i^{\text{new}}} \tilde{\ell}_{ij} \in \mathbb{R}^{B_r}$ .
  - 12:    Write  $\mathbf{O}_i \leftarrow \text{diag}(\ell_i^{\text{new}})^{-1} (\text{diag}(\ell_i) e^{m_i - m_i^{\text{new}}} \mathbf{O}_i + e^{\tilde{m}_{ij} - m_i^{\text{new}}} \tilde{\mathbf{P}}_{ij} \mathbf{V}_j)$  to HBM.
  - 13:    Write  $\ell_i \leftarrow \ell_i^{\text{new}}, m_i \leftarrow m_i^{\text{new}}$  to HBM.
  - 14:   **end for**
  - 15: **end for**
  - 16: Return  $\mathbf{O}$ .
-

# Reference

- [https://arxiv.org/pdf/2205.14135](https://arxiv.org/pdf/2205.14135.pdf)
- <https://www.cvmart.net/community/detail/7943>
- <https://www.youtube.com/watch?v=eMlx5fFNoYc>

# Lab6 - Objective

- Evaluate the performance of attention mechanisms by comparing:
  - The original PyTorch implementation
  - The FlashAttention v1 implementation
- Analyze the benefits of FlashAttention and explore its advantages over the standard approach.
- Conduct benchmarking with varying parameters and compare the results to gain deeper insights.

# Lab6 - Benchmark Script

- TA provide the benchmark script that does not require any modifications.
- Your task is to adjust only the parameters within the benchmark.
- The result will be outputted to a JSON file.
  - Execution time
  - FLOPs
  - Peak memory usage

```
benchmark_result.json
lab5 > benchmark_result.json > ...
1  {
2    "forward": {
3      "time(s)": 0.007095515976349513,
4      "FLOPS(TFLOPs/s)": 38.739664297876566
5    },
6    "backward": {
7      "time(s)": 0.018877355754375456,
8      "FLOPS(TFLOPs/s)": 36.40312638599925
9    },
10   "forward_backward": {
11     "time(s)": 0.025972871730724968,
12     "FLOPS(TFLOPs/s)": 37.04144402199095
13   },
14   "peak_memory_usage(MB)": 1288.00048828125
15 }
```

# Lab6 - Benchmark Script

- Test following parameters and compare the results.
  - `batch_size` : int
  - `seq_len (N)` : int
  - `num_heads` : int, (must be divisible by `emb_dim`)
  - `emb_dim (d)` : int
  - `impl` : str, (choose between Pytorch and Flash1)
  - `causal` : bool

# Lab6 - Benchmark Script

- [Colab download link](#) Connect to T4 GPU



```
# Install flash-attn Package (About 20 Min)
!pip install flash-attn==1.0.9 --no-build-isolation -q

Preparing metadata (setup.py) ... done
Building wheel for flash-attn (setup.py) ... done

import torch
!nvidia-smi
print("CUDA Usage :", torch.cuda.is_available())
print("GPU :", torch.cuda.get_device_name(0))

... Sun Nov  9 07:19:16 2025
+-----+-----+-----+
| NVIDIA-SMI 550.54.15      Driver Version: 550.54.15      CUDA Version: 12.4 |
| Persistence-M | Bus-Id      Disp.A    | Volatile Uncorr. ECC | | |
| GPU Name     | Pwr:Usage/Cap | Memory-Usage | GPU-Util  Compute M. |
| Fan Temp     | Perf          |             |          | MIG M. |
+-----+-----+-----+
| 0  Tesla T4      Off        00000000:00:04.0 Off | 0%       Default N/A |
| N/A 46C P8      11W / 70W | 0MiB / 15360MiB |          |          |
+-----+-----+-----+

Processes:
GPU GI CI PID Type Process name          GPU Memory Usage
ID ID
=====
No running processes found

CUDA Usage : True
GPU : Tesla T4
```

# Lab6 - Submission

- Plot the experimental data in a chart for better visualization.
- Analyze and explain your observations based on the collected data.
- Submit your report as a **lab6.pdf** file to eeclasse before **11/27 23:59**.

# HW4 - FlashAttention

- Spec
- Deadline : 12/7 23:59